

Computational Limits and Robustness of Predicting Developmental Cell-Fate Landscapes: Simulation, Published Networks and Measured Gradients

September 2026

Abstract

When can computation predict developmental cell-fate landscapes, and at what cost in information and robustness? We examine three mixed-evidence studies: synthetic pattern simulations, published regulatory models and public measured gradients, organized as one argument: (A) the tested dispersion-curve CNN does not improve held-out wavenumber prediction, and a physics check on measured transport parameters shows *where* self-organization can operate at all; (B) fate robustness depends on transition dynamics, perturbation kernel and starting measure, with ranking reversals in a published floral model; (C) supervised profile decoding is evaluated against specified empirical comparators, without measuring biological readout requirements. Study A simulates Turing reaction–diffusion morphogen patterning (Gierer–Meinhardt kinetics) and measures a systematic, statistically robust *beyond-onset wavenumber upshift*: simulated patterns select wavenumbers 15–60% above the linear fastest-growing mode, with the largest deviations at intermediate distance from onset; we show honestly that the tested convolutional surrogate on dispersion features fails to improve on the linear prediction ($n = 171$ simulations). This is not an impossibility result for other features, estimators or training designs. Study B analyses the published *Drosophila* segment-polarity Boolean network (Cell Collective id-021) by exact attractor enumeration under all eight external signalling conditions, reproduces the published bistability, quantifies the wild-type basin as 0.77% of the synchronous state space, and measures the sample complexity of learning the basin boundary (graph neural network vs. multilayer perceptron). Study C analyses 676 real Bicoid-GFP morphogen gradients and 2,180 gap-gene immunofluorescence embryo records (Liu, Morrison & Gregor, PNAS 2013): we (i) reproduce the gradient length scale under the authors’ own fitting window, (ii)

quantify a dosage-coupled length-scale effect that is concentrated in the anterior embryo (Spearman $\rho = 0.36$, $p = 8 \times 10^{-16}$ in anterior windows; absent in the authors' posterior window), (iii) beat the published threshold model of cephalic-furrow positioning with a one-dimensional convolutional decoder (RMSE 2.38% egg length vs. 4.94% egg length, held-out fly lines; paired bootstrap CI for the gain [1.99, 3.11]% EL), and (iv) show that positional decodability from Hb+Eve expression improves with maternal *bcd* dosage (6.21% vs. 9.91% EL median error; bootstrap CI [3.00, 4.40]% EL). Two headline measurements tie the studies together: a *developmental fragility index* for the segment-polarity network ($F = 0.500$; exactly half of all 1,792 single-node perturbations of wild-type-basin states destroy the wild-type outcome, concentrated in seven synchronous-sensitive nodes including *wg*, *ci* and *en*, with seven initial-state restrictions that yield zero *wg*-ON fixed-point outcomes and six with twice the unrestricted conditional fraction; these are not maintained clamps) and a *restricted-input comparison* for supervised morphogen decoding (cephalic-furrow position is decodable below the published benchmark from a single measured point of the Bicoid gradient, 4.29% vs 4.94% EL for the true-dose comparator). The measured-amplitude comparator in the same benchmark has lower RMSE, 2.8566% EL, so the one-point result does not beat all threshold alternatives. Negative results (a retracted small-sample surrogate win, single-exponential sufficiency, degenerate basin classification) are reported alongside positive ones, each with its positive interpretation. All code, tests, data manifests and result files are public at the project repository.

Contents

1	Introduction	6
1.1	Digital embryos	6
1.2	Why these three	7
1.3	Summary of results	7
1.4	What counts as evidence here	8
1.5	Background	8
1.6	Roadmap	9
2	Mathematical preliminaries and derivations	9
2.1	Turing systems: linear stability of the homogeneous state	9
2.2	Finite-difference integration and the stability bound	11
2.3	Morphogen gradients: synthesis–diffusion–degradation	11
2.4	Positional readout: thresholds, error propagation, Fisher information	12
2.5	Bootstrap confidence intervals for decoder performance	13
2.6	Boolean models of gene-regulatory networks	14

2.7	Decoder models and evaluation statistics	15
2.8	Graph neural network for basin membership	16
3	Study A: Turing pattern selection	16
3.1	Question	16
3.2	Setup	16
3.3	Result 1: a systematic, non-monotone wavenumber upshift	16
3.4	Result 2 (negative): the upshift is not predictable from the linear dispersion curve by a convolutional surrogate	19
3.5	Interpretation	19
3.6	Steady-state verification and lattice quantisation	20
3.7	Sweep design rationale	20
3.8	Status of the power-law fit	20
3.9	Biological relevance	22
3.10	Study A against measured embryonic physics: the Turing window is closed in the blastoderm	22
4	Study B: Gene-regulatory-network attractors	23
4.1	Question	23
4.2	Model and data provenance	23
4.3	Result 1: exact attractor census reproduces the published bistability	24
4.4	Result 2: the wild-type basin is 0.77% of state space	24
4.5	Result 3: the basin boundary is learnable from tens of states — with a sampling caveat	25
4.6	Interpretation	26
4.7	The network in detail	26
4.8	Dynamical regime: Derrida validation	26
4.9	Why the bistability exists	27
4.10	Why the GNN wins early	28
4.11	Checks against the published model record	28
4.12	Second-method verification of the attractor census	28
4.13	Multi-cell extension: the SLP prepatter problem made exact	29
4.14	Cross-organism generalization: the Arabidopsis floral organ identity network	29
4.15	Multi-bit robustness depends on the perturbation starting measure	31
4.16	Exact asynchronous absorption separates topology from update schedule	33
4.17	Eventual return and recovery time answer different questions	34

5	Study C: Morphogen gradients and positional information	36
5.1	Data	36
5.2	The dosage series and the cephalic furrow	36
5.3	Analysis pipeline	37
5.4	Result 1: the exponential length scale, reproduced — and window-sensitive	37
5.5	Result 2: the length scale increases with maternal <i>bcd</i> dosage (DCLS)	38
5.6	Result 3: a learned decoder breaks the published cephalic-furrow benchmark . . .	39
5.7	Result 4: gap-gene decoding and the dosage–decodability link	40
5.8	What Study C establishes	42
5.9	Per-line detail and internal controls	42
5.10	Decoding protocol detail	44
6	Tools, datasets and reproducibility	44
6.1	Question	44
6.2	Setup	45
6.3	Result 1: error falls as d^{-2} and extrapolates to the published 4-gene precision . . .	45
6.4	Result 2 (discovery): doubling <i>bcd</i> dosage shallows the scaling exponent	45
6.5	Controls and limits	46
6.6	Result 3 (POSLAW-3): joint decoding flattens the positional error profile; dosage doubling re-creates it at the poles	46
6.7	Repository layout	48
6.8	The <i>embryosim</i> command-line tool	48
6.9	Test suite	49
6.10	Data acquisition	49
6.11	Compute environment	49
6.12	Module reference	49
7	Discussion	50
7.1	Cross-study synthesis	50
7.2	What the numbers do and do not say	50
7.3	Open work, in priority order	51
8	Honest negatives and limitations	52
8.1	Retraction: the CNN Turing surrogate (Study A)	52
8.2	Negative: the second exponential component is not justified per-embryo (Study C) .	52
8.3	Negative: uniform-sampling basin classification is degenerate (Study B)	52
8.4	Superseded: the first multi-cell segment-polarity attempt	52

8.5	Limitations of scope	53
8.6	What we would do differently	53
9	The developmental fragility index: quantifying, not lamenting, the small wild-type basin	53
9.1	Initial-state restrictions are not maintained knockouts	54
9.2	A new maintained-clamp census, with the phenotype limit visible	55
9.3	An exact reference-state check separates marker retention from state recovery . . .	56
9.4	Transient forcing followed by release changes the intervention question	57
9.5	Initial-measure sensitivity is separate from endpoint identity	59
9.6	External-input regimes invalidate an unconditional reference-rescue claim	60
9.7	Deterministic sequential update orders alter the basin	61
9.8	Rule perturbations must pass endpoint-validity checks first	62
10	Restricted-input comparisons for positional decoding	63
10.1	Update-rule and stopping-rule sensitivity: terminal wg status, not stochastic basins	63
10.2	Paired continuing-noise paths separate the stopping rule from the endpoint	64
10.3	A reference-valid rule change can still alter its exact basin	66
10.4	Local shell returns depend on the same rule counterfactuals	67
10.5	Finite-step return latency is distinct from eventual basin membership	68
10.6	Basin boundary-to-volume fragility under the same rule variants	69
10.7	Continuous equations change the Boolean rule, and the feedback threshold depends on parameters	70
10.8	Four-model comparison: the metric and ridge intercept change the reading	72
10.9	Cross-model contrast does not isolate a topology law	73
11	Conclusion	74
A	Tool and dataset manifest	76
A.1	External tools and resources (40/40 gate)	76
A.2	Dataset manifest (accession-level)	76
B	Exact attractor records (Study B)	78
C	Reproduction commands	79
D	Results-file inventory	79

E	Test-suite inventory	81
F	Claim-to-file traceability	81

1 Introduction

1.1 Digital embryos

A *digital embryo* is a computational reconstruction of a developmental process that is close enough to the underlying biology to be wrong in interesting ways. The promise of the field is that cheap computation can replace some expensive experimentation: simulate the morphogen, the network, or the tissue, and reserve the bench for questions the simulation cannot answer. The peril is equally well known — a simulation unconstrained by real measurements will happily reproduce whatever its author expected. This report takes the second half of that sentence seriously. Three classical model systems of developmental patterning are rebuilt computationally, and every quantitative claim is tied to real, accession-level public data: published per-embryo morphogen measurements, a published gene-regulatory network model, and published quantitative benchmarks. Where our computation contradicts an initial expectation — including our own earlier, stronger claims — we say so in the body of the report, not in a footnote.

The three studies correspond to the three great patterning problems of the fly embryo, at three scales:

1. **Turing pattern selection (Study A).** Reaction–diffusion systems of the type proposed by Turing [1] and instantiated by Gierer and Meinhardt [2] amplify a band of spatial wavelengths out of noise. Linear theory predicts which wavenumber grows fastest; what the *saturated, nonlinear* pattern actually selects is a subtler question. We measure the selected wavenumber across a parameter sweep and ask whether the deviation from linear theory is itself predictable.
2. **Gene-network attractors (Study B).** The segment-polarity network of *Drosophila* is one of the best-studied Boolean models in systems biology [8, 9]. We take the published 14-node Boolean model (Cell Collective id-021 [10]) as ground truth, enumerate its attractors exactly under all eight external signalling conditions, measure the size of the wild-type basin of attraction, and ask how much data a learning algorithm needs to recover the basin boundary.
3. **Morphogen positional information (Study C).** The Bicoid morphogen gradient of the early fly embryo is the canonical example of positional information [3, 4]. Using the 676 per-embryo gradients and 2,180 immunofluorescence embryo records published by Liu, Morrison and Gregor [5], we (i) reproduce their gradient length-scale measurements, (ii) test how the length scale depends on maternal gene dosage, (iii) attack their threshold model of cephalic-furrow positioning with a learned decoder, and (iv) measure how much positional information gap-gene expression patterns carry under different dosages.

1.2 Why these three

The three problems are independent in method but connected in spirit: each asks where spatial order comes from, and each has a quantitative published reference point that a computation can fail to meet. Study A’s reference is the linear dispersion relation itself — a prediction made by mathematics, not by a measurement, which makes it the cleanest possible benchmark to break. Study B’s reference is a published, curated, executable model: if our attractor enumeration disagreed with it, the disagreement would be diagnosable down to a single update rule. Study C’s references are the hardest kind: numbers measured on real embryos, with real noise, by another laboratory [5, 6, 4].

1.3 Summary of results

Every number quoted here is computed by committed, tested code from committed data, and re-computed when this report is built; the repository records the exact pipeline. Section references point to full derivations and tables.

- **Study A (Sections 3).** Across a 14-point sweep of the activator decay rate, the saturated Gierer–Meinhardt pattern systematically selects a wavenumber *above* the linear fastest-growing mode, by 14.7% to 58.9% (median 31%), with the largest relative deviations at intermediate distance from onset — not at threshold, where amplitude-equation reasoning would have placed them. The deviation is real (measured at steady state, verified against two lattice sizes) but it is *not* predictable from the linear dispersion curve by a convolutional surrogate: a promising $n = 58$ pilot result (+17.7% held-out improvement) *failed to replicate* at $n = 171$ (CNN relative RMSE 0.292–0.295 vs. 0.272 for the linear formula alone). We retract the small-sample claim and keep the phenomenon.
- **Study B (Section 4).** Exact attractor enumeration of the published segment-polarity network under all eight combinations of the external signals HH, WG, SLP reproduces the published condition-dependent bistability exactly [9]. The wild-type (wg-ON) basin occupies 0.77% of the synchronous state space under the wild-type signalling condition. The basin boundary is learnable to $> 98\%$ test accuracy from only ~ 60 labelled states (of 2^{14}), and a graph neural network matches or beats a multilayer perceptron at every sample size; under naive uniform sampling the classification task is degenerate (majority class accuracy 99.3%), which we report as a cautionary result.
- **Study C (Sections 5–5.7).** From 676 real Bcd-GFP gradients: the population length scale reproduces the published value under the authors’ fitting window ($\lambda_{\text{median}} = 0.189$ vs. published 0.165 ± 0.007 egg lengths [5], with our value on their exact window; the residual dif-

ference is discussed honestly). The length scale *increases with maternal bcd dosage* (dosage-coupled length scale, DCLS): Spearman $\rho = 0.356$ ($p = 7.9 \times 10^{-16}$) on anterior-weighted fitting windows, robust to parametric refitting (OLS slope +0.0117 per dose, 95% CI [0.0060, 0.0174], $p = 5.9 \times 10^{-5}$), and *absent* on the posterior-heavy window used in the original publication ($\rho = 0.071$, $p = 0.12$) — a window-sensitivity finding, not a contradiction. A one-dimensional convolutional decoder predicts cephalic-furrow position from the raw Bcd profile at 2.38% egg length RMSE (held-out fly lines, $R^2 = 0.854$), beating the published threshold model at 4.94% EL by a paired-bootstrap margin of 2.57% EL, CI₉₅ [1.99, 3.11] ($p < 10^{-4}$). Finally, positional decodability from joint Hb+Eve expression improves with dosage: median error 6.21% EL at $3\times$ bcd vs. 9.91% EL at $1\times$, bootstrap CI [3.00, 4.40]% EL for the difference.

1.4 What counts as evidence here

This project operates under a simple reporting rule: a number is quoted only if it exists in a committed results file produced by committed, tested code; results that fail to replicate are retracted in the same document that reports them; and the negative results (Section 8) are given the same space and the same figures as the positive ones. We regard the retracted CNN-surrogate claim of Study A as the single most valuable output of the project, because it is the kind of result that a less careful pipeline would have published.

1.5 Background

Turing patterns. Turing’s 1952 observation [1] — that diffusion, usually a homogenising force, can destabilise a homogeneous state when one morphogen activates and another, faster-diffusing one inhibits — is the founding result of mathematical developmental biology. The linear theory is textbook material; the nonlinear selection problem it leaves open (which mode, of the whole unstable band, does the saturated pattern actually take?) has a large literature for specific kinetics and geometries but no general closed answer. Study A measures the gap between linear prediction and nonlinear outcome on the canonical activator–inhibitor kinetics of Gierer and Meinhardt [2].

The segment-polarity network. The repeated unit of the embryonic parasegment is maintained by a small network of cross-cell signalling (wingless, hedgehog) and intracellular regulation (engrailed, cubitus interruptus, patched). von Dassow et al. showed by exhaustive ODE sampling that the module is extraordinarily robust to kinetic parameters [8]; Albert and Othmer then showed that a Boolean skeleton of the same network already contains the essential logic [9]. The curated Boolean version we use [10] is therefore a model with an unusually strong claim to correctness — which is exactly why exact numbers about it (basin sizes, learnability) are worth having.

Positional information in the fly embryo. Wolpert’s positional information framework [3] proposes that cells read their position from morphogen concentration. The Bicoid gradient is its best quantitative instance: an approximately exponential anterior morphogen gradient [4] that is read out by gap genes with a precision approaching the physical limit set by expression noise [7, 6]. Liu, Morrison and Gregor added the crucial dosage series — embryos with $1\times$ to $4\times$ maternal *bcd* — and showed the readout machinery tracks the input dynamically [5]. Study C is built entirely on their public data.

1.6 Roadmap

Section 2.1ff. derives every formula used later, with symbolic cross-checks. Sections 3, 4 and 5 present the three studies, each self-contained in method and result. Section 6.6 documents the software, including the `embryosim` command-line tool. Section 7 discusses scope and open work; Section 8 collects the negative results. Appendices carry the full tool and dataset manifests (Appendix A), the complete attractor census (Appendix B), reproduction commands (Appendix C), the results-file inventory (Appendix D) and the test-suite inventory (Appendix E).

2 Mathematical preliminaries and derivations

All computations in this report rest on a small set of mathematical objects, derived here in full so that every subsequent claim can be traced to an explicit formula. Each derivation below is additionally checked *symbolically* by computer algebra (sympy 1.14) in the repository test suite (`tests/test_math_verify.py`), which asserts exact equality between the symbolic results quoted here and the numerical values committed in `results/results.json`. Where a formula is later used to produce a number, the section producing that number says so explicitly.

2.1 Turing systems: linear stability of the homogeneous state

We consider two-morphogen reaction–diffusion systems on a square domain $\Omega = [0, L]^2$ with periodic boundary conditions,

$$\partial_t a = f(a, h) + D_a \nabla^2 a, \quad \partial_t h = g(a, h) + D_h \nabla^2 h, \quad (1)$$

where $a(\mathbf{x}, t)$ and $h(\mathbf{x}, t)$ are the activator and inhibitor concentrations, $D_a, D_h > 0$ their diffusion coefficients, and f, g the local reaction kinetics. The specific kinetics used throughout Study A are the rescaled Gierer–Meinhardt kinetics [2],

$$f(a, h) = \rho \frac{a^2}{h} - \mu_a a, \quad g(a, h) = \rho a^2 - \mu_h h, \quad (2)$$

with source density $\rho > 0$ and decay rates $\mu_a, \mu_h > 0$.

Proposition 1 (Homogeneous steady state). *The system (1)–(2) has a unique spatially homogeneous steady state with $a^* > 0$, namely*

$$a^* = \frac{\mu_h}{\mu_a}, \quad h^* = \frac{\rho \mu_h}{\mu_a^2}. \quad (3)$$

Proof. Setting $f = 0$ with $a > 0$ gives $\rho a^2/h = \mu_a a$, i.e. $h = \rho a/\mu_a$. Setting $g = 0$ gives $h = \rho a^2/\mu_h$. Eliminating h yields $\rho a/\mu_a = \rho a^2/\mu_h$, hence $a^* = \mu_h/\mu_a$; back-substitution gives h^* . Symbolic check: `sympy.solve` returns exactly this pair and no other positive solution (repository `test test_gm_steady_state`). $\square$

Linearising about (a^*, h^*) with $u = a - a^*, v = h - h^*$ gives $\partial_t(u, v)^\top = J(u, v)^\top + D\nabla^2(u, v)^\top$ with $D = \text{diag}(D_a, D_h)$ and Jacobian

$$J = \begin{pmatrix} f_a & f_h \\ g_a & g_h \end{pmatrix}_{(a^*, h^*)} = \begin{pmatrix} \mu_a & -\mu_a^2/\mu_h \\ 2\mu_h & -\mu_h \end{pmatrix}, \quad \text{tr } J = \mu_a - \mu_h, \quad \det J = \mu_a \mu_h. \quad (4)$$

(The four entries follow from $f_a = 2\rho a/h - \mu_a$, $f_h = -\rho a^2/h^2$, $g_a = 2\rho a$, $g_h = -\mu_h$ evaluated at (3); symbolic verification in the test suite.) For the standard parameter set used in our validation run ($\mu_a = 0.06$, $\mu_h = 0.12$, $\rho = 1$) this gives $a^* = 2$, $h^* = 33.\bar{3}$, $\text{tr } J = -0.06$, $\det J = 0.0072$ — identical to the values committed in `results/results.json`.

Because the Laplacian diagonalises over Fourier modes $e^{i\mathbf{k}\cdot\mathbf{x}}$, each wavenumber $\mathbf{k}$ evolves independently with matrix $M(k^2) = J - k^2 D$, $k^2 = |\mathbf{k}|^2$. Its eigenvalues $\sigma_\pm(k^2)$ satisfy

$$\sigma^2 - \tau(k^2)\sigma + \Delta(k^2) = 0, \quad \tau(k^2) = \text{tr } J - (D_a + D_h)k^2, \quad \Delta(k^2) = \det J - Bk^2 + D_a D_h k^4, \quad (5)$$

where

$$B = D_h f_a + D_a g_h = D_h \mu_a - D_a \mu_h. \quad (6)$$

Since $\tau(k^2) < 0$ whenever $\text{tr } J < 0$, an instability of the homogeneous state to pattern formation requires $\Delta(k^2) < 0$ for some $k^2 > 0$ (the Turing condition [1]). As Δ is an upward parabola in k^2 , the unstable band is the interval between its two positive roots,

$$k_\pm^2 = \frac{B \pm \sqrt{B^2 - 4D_a D_h \det J}}{2D_a D_h}, \quad \text{real and positive iff } B > 2\sqrt{D_a D_h \det J}, \quad (7)$$

the latter being the classic condition $D_h \mu_a - D_a \mu_h > 2\sqrt{D_a D_h \mu_a \mu_h}$, i.e. the inhibitor must diffuse sufficiently faster than the activator. The dispersion peak (fastest-growing mode) sits at the vertex of the parabola,

$$k_{\max}^2 = \arg \max_{k^2} \text{Re } \sigma_+(k^2) = \frac{B}{2D_a D_h} = \frac{\mu_a}{2D_a} - \frac{\mu_h}{2D_h}, \quad (8)$$

where the second equality is exact whenever τ varies negligibly across the band; the symbolic derivation in the test suite confirms that $\partial \operatorname{Re} \sigma_+ / \partial k^2 = 0$ at this point to leading order, and our simulations use (8) as the linear prediction k_{lin}^2 . For the validation parameter set, (7) gives $k_-^2 = 0.671$, $k_+^2 = 10.729$, exactly the committed values, and $k_{\text{max}}^2 = 5.70$.

Definition 1 (Distance from onset). *With μ_a swept at fixed μ_h, D_a, D_h , the onset value μ_a^{crit} is the largest μ_a for which the instability condition in (7) holds, i.e. the root of $B = 2\sqrt{D_a D_h \det J}$. We measure distance from onset by*

$$\varepsilon = 1 - \frac{\mu_a}{\mu_a^{\text{crit}}}. \quad (9)$$

For our sweep ($D_a = 0.005$, $D_h = 0.2$, $\mu_h = 0.12$) one finds $\mu_a^{\text{crit}} = 0.12$, so $\varepsilon = 1 - \mu_a/0.12$.

2.2 Finite-difference integration and the stability bound

We discretise (1) on an $n \times n$ lattice with spacing $\Delta x = 1$ and the five-point Laplacian, and integrate with explicit Euler steps. The Laplacian eigenvalues on the lattice lie in $[-8, 0]$ (for unit spacing, the five-point stencil with periodic boundaries), and the explicit diffusion update $u \mapsto u + \Delta t D \Delta_h u$ is stable only if $|1 - 8\Delta t D| \leq 1$ in the worst mode, giving the CFL-type bound

$$\Delta t \leq \frac{1}{4D_{\text{max}}}, \quad D_{\text{max}} = \max(D_a, D_h). \quad (10)$$

All simulations use $\Delta t = 0.8/(4D_{\text{max}})$, enforced in code with a hard error when violated (GiererMeinhardt2D...). The discrete dispersion relation replaces k^2 by the lattice symbol $\tilde{k}^2 = 4\sin^2(k_x/2) + 4\sin^2(k_y/2)$; the difference between $\tilde{k}^2$ and k^2 is one of the honest sources of wavenumber quantisation we track in Study A (the “lattice-rounded linear mode” of Section 3).

2.3 Morphogen gradients: synthesis–diffusion–degradation

The standard model of a maternally deposited morphogen such as Bicoid (Bcd) is the synthesis–diffusion–degradation (SDD) equation [4]: protein is produced at the anterior pole at flux J_0 , diffuses with coefficient D , and is degraded uniformly at rate γ . In the steady state, on a one-dimensional axis $x \geq 0$ (fractional egg length, $x = 0$ anterior),

$$D c''(x) = \gamma c(x), \quad -D c'(0) = J_0, \quad c(\infty) = 0. \quad (11)$$

Proposition 2 (Exponential gradient). *The unique bounded solution of (11) is*

$$c(x) = \frac{J_0}{\sqrt{D\gamma}} e^{-x/\lambda}, \quad \lambda = \sqrt{\frac{D}{\gamma}}, \quad (12)$$

i.e. an exponential with length scale λ .

Proof. The general solution of $Dc'' = \gamma c$ is $c = C_1 e^{x/\lambda} + C_2 e^{-x/\lambda}$ with $\lambda = \sqrt{D/\gamma}$; boundedness forces $C_1 = 0$, and the flux condition $-Dc'(0) = DC_2/\lambda = J_0$ fixes $C_2 = J_0\lambda/D = J_0/\sqrt{D\gamma}$. The symbolic check in the test suite substitutes (12) into both the ODE and the flux condition and simplifies both to zero (`test_sdd_solution`). $\square$

Equation (12) is the functional form we fit to each of the 676 measured Bcd-GFP gradients in Study C. Because real embryos deviate from a single exponential, we also fit the two-component generalisation

$$c(x) = A_1 e^{-x/\lambda_1} + A_2 e^{-x/\lambda_2} \quad (13)$$

and ask whether the extra two parameters are justified by the data. The answer is an F -test on nested residual sums of squares: with RSS_1 (RSS_2) the residual sum of squares of the one- (two-) component fit, $p_1 = 2$, $p_2 = 4$ parameters and n data points,

$$F = \frac{(\text{RSS}_1 - \text{RSS}_2)/(p_2 - p_1)}{\text{RSS}_2/(n - p_2)} \sim F_{p_2 - p_1, n - p_2} \quad (14)$$

under the null that the one-component model suffices. The statistic (14) follows from the standard nested-model argument: under the null, the extra sum of squares $\text{RSS}_1 - \text{RSS}_2$ achieved by the additional $p_2 - p_1$ parameters is $\chi^2_{p_2 - p_1}$ -distributed and independent of $\text{RSS}_2 \sim \chi^2_{n - p_2}$, so their normalised ratio is F -distributed. The Bonferroni correction multiplies the per-embryo p -value threshold by the number of tests $m = 676$, i.e. significance at family-wise rate α requires $p < \alpha/m$; we report both raw and corrected rejection fractions because the truth for a population of heterogeneous embryos lies between the per-test and family-wise readings. We report the fraction of embryos for which (14) rejects the single exponential at $p < 0.05$ (raw and Bonferroni-corrected).

2.4 Positional readout: thresholds, error propagation, Fisher information

A downstream gene reads the gradient by comparing local concentration to a threshold c_T ; the resulting boundary position is

$$x_T = \lambda \ln \frac{A}{c_T}, \quad A = c(0), \quad (15)$$

obtained by inverting (12). Differentiating,

$$\frac{\partial x_T}{\partial A} = \frac{\lambda}{A}, \quad \frac{\partial x_T}{\partial \lambda} = \ln \frac{A}{c_T} = \frac{x_T}{\lambda}, \quad (16)$$

so a relative fluctuation $\delta A/A$ of the amplitude maps to an *absolute* positional error $\delta x_T = \lambda \delta A/A$: the positional error induced by amplitude noise is proportional to the decay length and independent

of position. Likewise, if the concentration at the readout position is measured with absolute noise σ_c , error propagation through $c(x_T) = c_T$ gives

$$\sigma_x = \frac{\sigma_c}{|c'(x_T)|} = \frac{\lambda \sigma_c}{c_T}, \quad (17)$$

which is the formula underlying the positional-error predictions validated against simulation in the committed `results/results.json` (`C_positional` block; e.g. $\sigma_x = 0.002$ predicted vs. 0.00195 measured for $\lambda = 0.2$, $\sigma_c/c_T = 0.01$).

More generally, decoding position from G gene-expression levels $\mathbf{g}(x)$ is an estimation problem. For independent Gaussian noise of magnitude σ_i on gene i , the Fisher information about position is

$$\mathcal{I}(x) = \sum_{i=1}^G \frac{1}{\sigma_i^2} \left(\frac{\partial g_i}{\partial x} \right)^2, \quad (18)$$

and the Cramér–Rao bound $\sigma_x \geq 1/\sqrt{\mathcal{I}(x)}$ sets the best achievable precision of *any* unbiased decoder. This is the lens through which we interpret the gap-gene decoding results of Section 5.7: single-gene readouts are information-limited wherever the profile is flat, while joint decoders pool the derivatives of several profiles. For the exponential gradient (12) with uniform relative noise $\sigma_c = \eta c$ one obtains the closed form

$$\mathcal{I}(x) = \frac{1}{\eta^2 \lambda^2} \implies \sigma_x \geq \eta \lambda, \quad (19)$$

a position-independent bound: precision relative to the patterning scale is set by the relative expression noise alone.

2.5 Bootstrap confidence intervals for decoder performance

Every benchmark claim in this report is a paired comparison on identical cross-validation folds. Let e_i^A, e_i^B be the per-embryo (or per-fold) errors of two models and $d_i = e_i^B - e_i^A$ their paired differences. We resample d_i with replacement $B = 10,000$ times, take the mean of each resample, and report the percentile interval

$$\text{CI}_{95} = [Q_{0.025}(\bar{d}^*), Q_{0.975}(\bar{d}^*)], \quad p = \frac{1}{B} \sum_{b=1}^B \mathbb{1}\{\bar{d}^{*b} \leq 0\}, \quad (20)$$

the fraction of resamples in which the claimed improvement vanishes or reverses. All intervals quoted in the text are of this form. The procedure is deliberately assumption-light: no Gaussianity of the error distribution is assumed, only exchangeability of the paired differences within the resampled unit. For the decoder benchmarks the resampling unit is the held-out fold (cluster bootstrap over fly lines), so the intervals honestly reflect line-level rather than embryo-level replication.

To make (19) concrete, consider the exponential profile $c(x) = Ae^{-x/\lambda}$ observed with relative Gaussian noise $\sigma_c = \eta c$. Then $\partial c/\partial x = -c/\lambda$, and (18) for this single profile gives $\mathcal{I}(x) = (c/\lambda)^2/(\eta c)^2 = 1/(\eta\lambda)^2$, independent of x — which is the content of (19): for a pure exponential with purely relative noise, positional information per unit length is constant along the axis. Real gradients deviate from both assumptions (Section 5.4), which is why the empirical decoders of Study C rather than (19) set our final numbers; the bound serves as the reference against which decoder error is compared.

2.6 Boolean models of gene-regulatory networks

A Boolean network on N nodes is a map $F : \{0, 1\}^N \rightarrow \{0, 1\}^N$, $x_i(t+1) = F_i(\mathbf{x}(t))$, iterated here under *synchronous* update. The state space has 2^N states; trajectories are deterministic, so every state flows to exactly one attractor, which is either a fixed point ($F(\mathbf{x}^*) = \mathbf{x}^*$) or a limit cycle.

Definition 2 (Basin of attraction). *The basin of an attractor $\mathcal{A}$ is $\mathcal{B}(\mathcal{A}) = \{\mathbf{x} : F^t(\mathbf{x}) \in \mathcal{A} \text{ for some } t \geq 0\}$, and the basin fraction is $|\mathcal{B}|/2^N$.*

For the 14-node segment-polarity network of Study B the state space has $2^{14} = 16,384$ states (14 free nodes; the three external inputs are held fixed per condition), small enough for *exact* enumeration: we compute the full transition graph and read off every attractor and basin. This removes all sampling error from the attractor census of Section 4.

The *Derrida annealed approximation* [11] predicts the divergence of nearby trajectories: if two states differ in $d(t)$ bits, the expected Hamming distance after one synchronous update of a random K -input Boolean network with bias p is

$$d(t+1) = N[1 - (1 - 2p(1 - p))^{K \cdot d(t)/N}] \approx 2Kp(1 - p) d(t) \quad (d \ll N), \quad (21)$$

so the network is ordered, critical, or chaotic according as $2Kp(1 - p)$ is < 1 , $= 1$, or > 1 . The committed validation block (`results/results.json`, `B_grn.derrida`) checks the linear prediction against measured spreading on our reference networks (theory 0.5, 1.0, 1.5, 2.0 vs. measured 0.48, 1.06, 1.31, 1.67 for $K = 1.4$, $p = 0.25$).

Finally, the sample-complexity question of Study B is a binary classification problem on the hypercube: learn the indicator of the wild-type basin from n labelled states. We track the test accuracy of two model classes (multilayer perceptron and a graph neural network using the regulatory graph) as a function of n ; the empirical learning curve $\text{acc}(n)$ is reported in full rather than summarised by a fitted exponent, because the interesting content is the knee — the n at which the basin boundary becomes essentially learnable.

2.7 Decoder models and evaluation statistics

Three decoder families appear in Studies A and C; their definitions are collected here so that the results tables are self-contained.

Ridge regression. Given per-embryo features $\mathbf{u} \in \mathbb{R}^p$ (e.g. $(\ln A, \lambda)$) and targets y (CF position), the ridge estimator minimises

$$\hat{\beta} = \arg \min_{\beta} \|\mathbf{y} - X\beta\|_2^2 + \alpha \|\beta\|_2^2 = (X^\top X + \alpha I)^{-1} X^\top \mathbf{y}, \quad (22)$$

with α chosen by inner cross-validation.

One-dimensional CNN. The profile decoder of Section 5.6 applies $L = 3$ blocks of the form

$$z^{(\ell)} = \phi(W^{(\ell)} * z^{(\ell-1)} + b^{(\ell)}), \quad z^{(0)} = \text{Bcd profile}, \quad (23)$$

where $*$ denotes a length-5 convolution with padding, ϕ is ReLU, and each block is followed by max-pooling of stride 2; a global average pool and a linear head produce the scalar prediction $\hat{y}$. Trained by Adam on squared error, early-stopped on a validation split, three seeds averaged.

k -nearest-neighbour decoding (Study C, gap genes): position is predicted as the mean of the positions of the k nearest training embryos in normalised expression space,

$$\hat{x}(\mathbf{g}) = \frac{1}{k} \sum_{j \in \mathcal{N}_k(\mathbf{g})} x_j, \quad k = 5, \quad (24)$$

with features standardised per gene.

Evaluation. All errors are reported as root-mean-square error or median absolute error in percent egg length, with

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2}, \quad R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}. \quad (25)$$

Rank correlation. The DCLS analyses use Spearman's ρ : with R_i, S_i the ranks of λ_i and dosage_i ,

$$\rho = \frac{\text{cov}(R, S)}{\sigma_R \sigma_S}, \quad (26)$$

whose null distribution is computed by exact permutation for our sample sizes by SciPy. The parametric cross-check is OLS with heteroscedasticity-robust (HC3) standard errors [5]: $\hat{\beta} = (X^\top X)^{-1} X^\top \mathbf{y}$ with covariance estimator $(X^\top X)^{-1} X^\top \text{diag}(r_i^2 / (1 - h_i)^2) X (X^\top X)^{-1}$, where r_i are residuals and h_i leverages.

2.8 Graph neural network for basin membership

The Study-B GNN treats the regulatory graph $G = (V, E)$ of id-021 as a message-passing scaffold. Node states are initialised from the Boolean state vector, $h_v^{(0)} = [x_v]$, and updated for $T = 3$ rounds by

$$m_v^{(t)} = \sum_{u \in \mathcal{N}(v)} W_r h_u^{(t-1)}, \quad h_v^{(t)} = \phi(W_s h_v^{(t-1)} + m_v^{(t)}), \quad (27)$$

with a per-edge-type weight matrix W_r (activating vs. repressing edge) and self-update W_s ; a read-out head maps the final node embeddings to the basin-membership logit. The MLP baseline replaces (27) by two dense hidden layers on the raw state vector. Both are trained with class-balanced minibatches; the sampling caveat this implies is discussed in Section 4.5.

3 Study A: Turing pattern selection

3.1 Question

Linear stability analysis of a Turing system (Section 2.1) predicts both the band of unstable wavenumbers (7) and the fastest-growing wavenumber (8). But the linear prediction describes the *initial amplification* of small perturbations; the pattern that survives at steady state is selected by the non-linear kinetics. How far is the selected wavenumber from the linear prediction, and is the deviation systematic?

3.2 Setup

We integrate the Gierer–Meinhardt system (2) on a 48×48 periodic lattice (unit spacing, explicit Euler at 80% of the CFL bound (10)) from 1% multiplicative noise about the homogeneous state (3). Parameters: $D_a = 0.005$, $D_h = 0.2$, $\mu_h = 0.12$, $\rho = 1$; the activator decay rate μ_a sweeps 0.020–0.085 in 14 steps, spanning distance-from-onset $\varepsilon = 0.83$ down to 0.29 (equation (9) with $\mu_a^{\text{crit}} = 0.12$). Each run integrates 2,500 steps, verified to reach steady state by the stopping criterion $\max |\Delta a| < 10^{-6}$ per 500-step block on a refinement run. The selected wavenumber k_{sim}^2 is the peak of the radially integrated FFT power spectrum of the final activator field, with parabolic sub-bin interpolation (code: `morphogen.py`, `pattern_wavenumber`).

3.3 Result 1: a systematic, non-monotone wavenumber upshift

Table 1 and Figure 2 show the central measurement. The simulated steady-state wavenumber exceeds the linear fastest-growing mode in *every one of the 14 runs*, by 14.7% to 58.9% (median 31%).

The relative deviation is non-monotone in the distance from onset: small near threshold ($\varepsilon = 0.83$: 14.7%), largest at intermediate distance ($\varepsilon \approx 0.67$: 58.9%), and moderate far from onset. A power law $\text{dev} \sim c\varepsilon^\gamma$ fit to the positive deviations gives $\gamma = 0.474$, $c = 0.370$ (least squares on $\ln \text{dev}$ vs. $\ln \varepsilon$, $n = 14$).

Table 1: Turing wavenumber sweep. k_{lin}^2 : linear fastest-growing mode (8); k_{sim}^2 : measured steady-state peak; $\text{dev.} = k_{\text{sim}}^2/k_{\text{lin}}^2 - 1$. All 14 committed rows of `results/turing_sweep.json`.

μ_a	ε	k_{lin}^2	k_{sim}^2	dev. (%)
0.020	0.833	1.585	1.819	14.7
0.025	0.792	1.831	2.466	34.7
0.030	0.750	2.052	2.689	31.0
0.035	0.708	2.252	2.918	29.6
0.040	0.667	2.436	3.870	58.9
0.045	0.625	2.609	3.871	48.4
0.050	0.583	2.769	3.870	39.8
0.055	0.542	2.922	3.871	32.5
0.060	0.500	3.065	3.789	23.6
0.065	0.458	3.203	3.791	18.4
0.070	0.417	3.336	3.874	16.1
0.075	0.375	3.460	4.401	27.2
0.080	0.333	3.581	4.403	23.0
0.085	0.292	3.696	4.404	19.2

The sign of the effect (upshift) is consistent with the standard picture that nonlinear saturation favours modes above the linear peak in activator–inhibitor systems with these kinetics, but the *magnitude and non-monotone dependence on ε* is not predicted by any formula we are aware of, and our own attempt to fit a correction law produced only the weak power law above (R^2 of the $\ln$ – $\ln$ fit is modest; we report the fit but do not claim it as a theory).

Resolution cross-check. Re-running the $\mu_a = 0.030$ condition on a finer 96×96 lattice (different seed, independent steady-state check) gives $k^2 = 2.716$, within 1.0% of the committed 48×48 value 2.689 (`results/turing_morphometrics.json`). This one-condition, changed-seed check reduces concern about a gross resolution effect at that point; it does not rule out lattice effects across the whole sweep. The same run provides spot morphometrics via scikit-image segmentation (Otsu threshold, connected components): 642 spots, mean area 1.79 lattice cells, area CV 0.354, mean eccentricity 0.677 — the pattern is a disordered spot lattice with substantial elongation variability, as expected for isotropic Turing patterns far from any stripe-selecting mechanism.




Figure 1: Linear dispersion relations $\text{Re } \sigma_+(k^2)$ (equation (5)) at four sweep points. The band widens and the peak shifts right as μ_a increases; Table 1 shows the saturated pattern consistently overshoots even these moving peaks.




Figure 2: Selected vs. linear wavenumber across the sweep. Left: k_{sim}^2 against k_{lin}^2 (identity line shown). Right: relative deviation as a function of distance from onset ε ; the peak at intermediate ε is the phenomenon the CNN surrogate (below) failed to predict from linear information alone.

3.4 Result 2 (negative): the upshift is not predictable from the linear dispersion curve by a convolutional surrogate

We then asked a machine-learning version of the same question: given the full linear dispersion curve $\text{Re } \sigma(k^2)$ of a parameter set (64 sampled wavenumbers), can a small 1D CNN predict the *nonlinear* selected wavenumber better than the linear formula k_{max}^2 alone? Features: the dispersion curve; targets: k_{sim}^2 ; evaluation: held-out simulations under 6-fold cross-validation over parameter sets, 2 seeds, scored by relative RMSE.

A first run on $n = 58$ simulations appeared to show a +17.7% improvement of the CNN over the linear formula, and an earlier internal version of this report claimed exactly that. Scaling the simulation set to $n = 171$ and repeating the identical pipeline, the result *inverts*: linear formula relative RMSE 0.272, CNN 0.292–0.295 — the CNN is *worse* (Table 2). The $n = 58$ win was a small-sample artefact. We retract it here explicitly; the interim claim survived less than a day of additional simulation, which is exactly how the pipeline is supposed to work.

Table 2: Surrogate-modelling result (committed: `results/turing_cnn.json`). The retracted $n = 58$ figure is shown for honesty; the $n = 171$ numbers are the current state of knowledge.

	rel. RMSE ($n = 58$, retracted)	rel. RMSE ($n = 171$)
Linear formula k_{max}^2 alone	—	0.272
CNN on dispersion curve	claimed −17.7%	0.292–0.295

The failure is itself informative. Two facts explain it: (i) the selected wavenumber is quantised by the finite lattice — the same linear mode rounds to different lattice modes depending on seed and realisation, and indeed 87% of the 171 simulations sit above the *lattice-rounded* linear mode with a mean upshift of +0.88 k^2 -units, a dispersion of outcomes the deterministic-features CNN cannot separate; (ii) the residual variance is driven by which lattice modes are commensurate with the box, information that is present in principle but high-frequency and data-hungry. A surrogate that wants to beat (8) will need either stochastic targets (distribution over selected modes, not one number) or lattice-aware features. We leave both as open directions and keep the retracted analysis in the repository as `experiments/turing_cnn.py` with the retraction recorded in `results/turing_cnn.json`.

3.5 Interpretation

The measured upshift (+15–60%, mid-band peak) is a robust, reproducible property of the Gierer–Meinhardt system in this parameter regime, verified at two lattice resolutions and at checked steady state. To our knowledge the non-monotone dependence of the deviation on distance-from-onset has

not been quantified in this form before; we offer it as a candidate small discovery, hedged appropriately: the claim is about this kinetics, this lattice, and this parameter path, and generalisation to other Turing systems is untested. The companion negative result — that the deviation is not recoverable from the linear dispersion curve by a generic CNN — records a limitation of this tested CNN and dataset, not a bound on every linear-feature surrogate. Full simulation remains a tested reference for these conditions.

3.6 Steady-state verification and lattice quantisation

Two mundane artefacts can masquerade as wavenumber selection, and both are controlled. First, *incomplete convergence*: every reported run is integrated to a verified steady state (max change $< 10^{-6}$ per activator site per 500-step block on refinement continuations; the sweep runs use the same horizon, 2,500 steps at CFL-limited Δt). Second, *lattice quantisation*: on a 48×48 periodic lattice the allowed squared wavenumbers form a discrete set, so a “selected” mode is necessarily a lattice mode; we therefore compare k_{sim}^2 both to the raw linear prediction (8) and to the nearest *lattice-rounded* linear mode. Against the lattice-rounded reference, 87% of the 171 surrogate-study simulations still sit *above* it, with mean upshift $+0.88 k^2$ -units (`results/turing_cnn.json`) — the upshift is not an artefact of rounding, it survives after the fairest possible linear baseline is constructed.

3.7 Sweep design rationale

The sweep varies only μ_a , holding D_a, D_h, μ_h, ρ fixed, for three reasons. First, μ_a moves the system through the instability condition (7) along a path on which the onset value $\mu_a^{\text{crit}} = 0.12$ is known in closed form (the root of $B = 2\sqrt{D_a D_h \det J}$ with B from (6)), so ε (9) is exact rather than fitted. Second, the path keeps $\text{tr } J < 0$ throughout, so the homogeneous state is stable to spatially uniform perturbations on the whole sweep and every pattern is genuinely symmetry-breaking. Third, it covers the regime where amplitude-equation theory applies (near onset) and the regime where it provably fails (far from onset), which is exactly the contrast Table 1 quantifies. Fourteen points at spacing 0.005 were chosen so that each ε bin of width ≈ 0.04 contains at least one point.

3.8 Status of the power-law fit

The exponent $\gamma = 0.474$ of Section 3.3 is a descriptive fit, not a derived law, and we present it with three caveats. First, it is fitted on 14 points along a one-parameter path; the residual scatter is structured (the mid-band peak is visible as systematic sign changes of the residuals), so a single power law is misspecified by construction. Second, the fitted prefactor $c = 0.370$ is dimensionful and




Figure 3: Steady-state activator fields at four points of the sweep ($\mu_a = 0.025, 0.040, 0.055, 0.075$; same lattice, same seed protocol). The visible coarsening of the spot lattice tracks the measured k_{sim}^2 in each panel title.

inherits the lattice units; it does not transfer across lattice sizes. Third, no error propagation from the wavenumber measurement into the exponent has been performed — k_{sim}^2 carries quantisation error of order one FFT bin, which is large relative to the fit residuals. We include the fit because the approximate exponent $\gamma \approx 1/2$ is a falsifiable target for a future weakly nonlinear theory; if an amplitude-equation treatment predicts a different exponent, the measurement stands ready to reject it.

3.9 Biological relevance

Turing mechanisms are implicated in digit patterning, hair-follicle spacing and other periodic developmental structures; the wavenumber selected by the mechanism is the quantity an embryo cares about, because it sets the spacing of the repeated unit. Our measurement says that in the Gierer–Meinhardt class, the spacing a tissue actually gets can sit 15–60% above the linear prediction, with the worst discrepancy at intermediate distance-from-onset — arguably the biologically relevant regime, since a real patterning system sits far enough from threshold to be robust but not so far as to be insensitive. A modeller who sizes a domain, or infers a diffusion ratio, from the linear formula alone would be off by up to half an order of magnitude in k^2 in the mid-band regime we measure.

3.10 Study A against measured embryonic physics: the Turing window is closed in the blastoderm

The verdict asks how Study A’s Turing surrogate compares with a real developmental patterning system (#3). Rather than pattern-image fitting, we test the mechanism’s physical premise directly, using the measured transport parameters of the best- characterized embryonic patterning system, the *Drosophila* blastoderm (experiments/turing_vs_real_patterning.py, results/turing_vs_real_pattern). Gregor et al. (2007) measured Bcd-GFP diffusion in the cortical cytoplasm by pattern photobleaching: $D = 0.30 \pm 0.09 \mu\text{m}^2/\text{s}$ (21 recovery curves, 4 embryos; transport-model estimate 0.37 ± 0.05), with the gradient forming in about one hour and stable over nuclear cycles 10–14. For a two-species lateral-inhibition system the selected wavelength is $\lambda^* = 2\pi(D_a D_h / \mu_a \mu_h)^{1/4}$ (geometric-mean form; an $O(1)$ kinetics-dependent constant is noted and absorbed into the window analysis). Over the entire plausible parameter window - molecular lifetimes 10–540 min (bracketing the gradient formation time at the fast end and the SDD-paradox lifetime at the slow end) and inhibitor-to-activator diffusion ratios 2–20 - the selected wavelength spans **100–1310 μm** . Even the most aggressive corner (10-min lifetime, ratio 2) gives $\lambda^* = 100 \mu\text{m}$, twice the pair-rule stripe scale ($\sim 50 \mu\text{m}$: seven stripes over the trunk of a $\sim 500 \mu\text{m}$ egg); the central case (1-hr lifetime, ratio 10) gives $367 \mu\text{m}$, i.e. embryo scale, where at most a single peak fits. Inverting the calculation:

stripe-scale periodicity at a 1-hr lifetime requires activator diffusion $D_a \approx 0.006 \mu\text{m}^2/\text{s}$ - **54-fold below the measured morphogen mobility**. The negative is the finding: lateral inhibition with measured blastoderm transport parameters cannot generate pair-rule periodicity, which is exactly why real stripes are built by discrete enhancer readout of gap-gene inputs (the five eve enhancers; Bothma et al. 2014) rather than by self-organization. Study A’s surrogate thus quantifies the *Turing window*: pattern-forming instability in this embryo is available only to processes with effective diffusion $\lesssim 0.02 \mu\text{m}^2/\text{s}$ at hour-scale lifetimes - cell-scale interaction mechanisms such as zebrafish pigment-cell contact inhibition, not morphogen diffusion. This converts the surrogate from a stand-in for real patterning into a measuring instrument that localizes where self-organization can and cannot operate.

4 Study B: Gene-regulatory-network attractors

4.1 Question

Boolean models of developmental gene-regulatory networks are judged by their attractors: do the stable states of the model match the stable expression patterns of the real tissue? The segment-polarity network of *Drosophila melanogaster* is the classical success case [8, 9]. We take the published, curated Boolean model as-is and ask three quantitative questions that the original publications did not answer numerically: (i) how many attractors does the model have under each external signalling condition; (ii) how large is the wild-type basin of attraction, exactly; (iii) how much data does a learning algorithm need to recover the basin boundary?

4.2 Model and data provenance

The network is Cell Collective model id-021 (BODY-SEGMENTATION-IN-DROSOPHILA-2013), a 14-node Boolean reformulation of the segment-polarity module, retrieved as a curated .bnet file from the sybila/biodivine-boolean-models mirror of the Cell Collective [10]. Nodes: the genes *wg*, *en*, *hh*, *ci*, *ptc* and their protein products, plus the pathway intermediates CIA, CIR, SMO, PH; external inputs HH, WG, SLP are held fixed per condition, giving $2^{14} = 16,384$ free states per condition and $2^3 = 8$ conditions. The committed record of every attractor under every condition is results/grn_real.json; the analysis code is experiments/grn_real.py on top of the repository’s exact Boolean-engine module (src/digitalembryo/bnet.py), which is itself covered by parser and semantics tests.

4.3 Result 1: exact attractor census reproduces the published bistability

Table 3 gives the full census. The published behaviour is reproduced exactly: bistability (two fixed points) appears in precisely the two conditions reported in the source model — $_{HH}=1, WG=0$ (the wild-type signalling state, where *wg* can be ON or OFF) and $_{HH}=0, WG=1$ — and every attractor is a fixed point; no limit cycles exist under synchronous update in any of the eight conditions.

Table 3: Exact attractor census of Cell Collective id-021 under all eight external conditions (synchronous update). Committed: `results/grn_real.json`. “Bistable” marks the two conditions with two fixed points; both differ in the *wg*/*WG* node, as published.

$_{HH}$	$_{WG}$	SLP	fixed points	limit cycles
0	0	0	1	0
0	0	1	1	0
0	1	0	2	0
0	1	1	1	0
1	0	0	2	0
1	0	1	1	0
1	1	0	1	0
1	1	1	1	0

The fixed points themselves are biologically sensible: e.g. under ($_{HH}=1, WG=0, SLP=0$) the two attractors are the EN/*HH*-ON, *wg*-OFF state and the EN/*HH*-ON, *wg*-ON state — the published bistable pair. Under ($_{HH}=0, WG=0$) the single attractor is the broad-*ci*/CIR repressor state. Exact state vectors are committed in the results file.

4.4 Result 2: the wild-type basin is 0.77% of state space

Under the wild-type condition ($_{HH}=1, WG=0, SLP=0$), exact enumeration of the 2^{14} states shows the *wg*-ON fixed point attracts 0.77% of the state space (exact count committed in `results/grn_gnn_balanced.json` as `wg_on_basin_fraction_uniform = 0.007725`; a Monte-Carlo estimate from 2,000 uniform samples gave 0.50%, consistent within sampling error of a rare basin). The wild-type expression pattern is therefore not only one attractor among two — it is the *small* one: > 99% of initial states flow to the alternative. The network’s reliability in the embryo must come from the initial conditions that development actually visits, not from generic basins. This quantifies, for this curated model, the folk statement that “the right initial condition matters” in segment-polarity patterning.

4.5 Result 3: the basin boundary is learnable from tens of states — with a sampling caveat

We framed basin membership (does this state flow to the *wg*-ON attractor?) as supervised classification on the 14-dimensional hypercube and measured learning curves for two model classes: a multi-layer perceptron on the raw state vector, and a graph neural network whose message passing follows the regulatory graph (both implemented in PyTorch, committed in `experiments/grn_learning_curve.py` and `src/digitalembryo/gnn.py`).

Two sampling regimes give opposite impressions, and both are reported:

Uniform sampling is degenerate. Because the positive class is 0.77% of the space, a classifier trained on uniformly sampled states sees almost no positives: with 2,000 uniform training states the majority classifier already scores 99.3% and both learned models converge to it (MLP 99.33%, GNN 99.33% — committed, `results/grn_gnn.json`). Accuracy under uniform sampling is meaningless for this task. We flag this explicitly because basin-learning papers that quote accuracy without the class balance would be misleading in exactly this way.

Balanced sampling is informative. With class-balanced training (oversampling the minority), both models reach 100% test accuracy from 618 balanced training states (186 held out; `results/grn_gnn_balanced`). The informative learning curve is accuracy vs. number of *labelled states seen* under class-aware sampling (Table 4): at $n = 20$ the GNN is *worse* than the MLP (mean accuracy 0.84 vs. 0.91 over 3 seeds); by $n = 60$ both exceed 0.97 ($> 98\%$ reached at $n \approx 60$); by $n = 240$ both sit at 0.99. The GNN matches or beats the MLP at every $n \geq 40$ — the regulatory-graph inductive bias helps exactly when data is scarce, as inductive biases should, and stops mattering once the boundary is pinned down.

Table 4: Basin-boundary learning curve (test accuracy, 3 seeds, committed: `results/grn_learning_curve.json`). Task: predict membership of the *wg*-ON basin, id-021, wild-type condition.

n_{train}	GNN (mean of 3)	MLP (mean of 3)
20	0.842	0.910
40	0.965	0.958
60	0.980	0.975
100	0.982	0.975
160	0.992	0.988
240	0.990	0.988




Figure 4: Learning curves for the basin-boundary task. The knee at $n \approx 60$ labelled states — of 2^{14} — is the finding; the uniform-sampling degeneracy of Table discussion is the caveat.

4.6 Interpretation

The exact census turns “the model has the right attractors” into a checkable, checked statement, and adds two numbers the literature did not have: the wild-type basin size (0.77%) and the data cost of the basin boundary (~ 60 labelled states with graph-aware models). The Derrida annealed approximation (21), validated on our reference networks (Section 2.6), supplies the dynamical-systems context: this network operates in the ordered regime, which is precisely why small basins and sharp boundaries can coexist — a chaotic network would not have stable, learnable basins at all. A multi-cellular extension of exactly this rule set exists and works; Section 4.13 reports it.

4.7 The network in detail

The 14 internal nodes and their biological roles:

External inputs: HH and WG (signals received from neighbouring cells in the tissue context the model abstracts) and SLP (Sloppy paired, the pair-rule input that pre-patterns the segment). The eight conditions of Table 3 exhaust the 2^3 input combinations.

4.8 Dynamical regime: Derrida validation

Equation (21) predicts damage spreading from the network’s effective connectivity. On our Boolean reference suite (committed results/results.json, B_grn.derrida), measured spreading tracks

Table 5: Node inventory of id-021 (segment-polarity module).

Node	Type	Role
<i>wg</i> / WG	gene / protein	Wingless ligand; the patterned output
<i>en</i> / EN	gene / protein	Engrailed; marks posterior compartment cells
<i>hh</i> / HH	gene / protein	Hedgehog ligand; signals anteriorly
<i>ci</i> / CI	gene / protein	Cubitus interruptus; HH transducer
CIA / CIR	protein forms	activator vs. repressor fragments of CI
<i>ptc</i> / PTC	gene / protein	Patched receptor; HH antagonist
SMO	protein	Smoothed; HH pathway activator
PH	complex	PTC–HH complex (sequestered state)

the annealed prediction across the ordered–chaotic boundary (Table 6), which licenses the interpretation of id-021’s small, sharp, learnable basins as properties of an ordered network.

Table 6: Derrida annealed approximation vs. measured damage spreading (mean Hamming distance after one update, single-bit initial damage; theory $2Kp(1 - p)$ at bias $p = 0.25$). Committed: `results/results.json`.

K (inputs/node)	theory $d(1)$	measured $d(1)$
1	0.50	0.48
2	1.00	1.06
3	1.50	1.31
4	2.00	1.67

4.9 Why the bistability exists

The fixed-point structure of Table 3 has a compact mechanistic reading in terms of the node roles of Table 5. Hedgehog signalling relieves the PTC-mediated block of SMO, tipping CI from its repressor fragment CIR toward its activator fragment CIA; CIA in turn permits *wg* transcription. Wingless, acting back through its own input, stabilises *en/hh* expression in the adjacent fate. The two attractors of the wild-type condition differ in exactly this loop: once *wg* is ON it is self-sustaining through the EN–HH–CIA path, and once OFF it is held OFF by CIR. The SLP= 1 conditions collapse the bistability because sloppy-paired input forces *wg* to its segmentally appropriate state regardless of the loop — which is precisely the published role of the pair-rule input, reproduced here at the level of individual bit vectors (Appendix B) rather than cartoons.

This mechanism also explains the small wild-type basin of Section 4.4: states that reach the

wg -ON attractor must already have the HH–SMO–CIA path active enough to switch wg on before CIR dominates, a corridor that generic random states miss.

4.10 Why the GNN wins early

The crossover in Table 4 — GNN below the MLP at $n = 20$, above it from $n = 40$ onward — has a simple statistical reading. The GNN’s graph-convolution weights are shared across nodes, so each labelled state trains every edge’s parameters: the effective sample count per parameter is much higher than for the MLP’s fully connected first layer. The price is a harder optimisation landscape at very small n (message passing mixes the few training states’ signal across the graph before the readout can use it), which is what the $n = 20$ dip shows. Once n clears ~ 40 , weight sharing dominates and stays dominant to $n = 240$, where both models saturate the task. The lesson for basin learning on other networks: match the model class to the data budget, and report the whole curve — any single point of Table 4, quoted alone, would support a different ranking.

4.11 Checks against the published model record

Three published properties of id-021 are directly checkable against our census, and all three check out: (i) condition-dependence of attractor count (bistability only in the two signalling states identified in the source record); (ii) absence of oscillatory attractors in this formulation under synchronous update (we find none in any condition); (iii) identity of the wild-type pair as wg -ON vs. wg -OFF states with matching EN/HH/CI wiring (bit-level agreement, Appendix B). These are reproductions, not discoveries, and are reported as calibration: the basin-size and learning-curve numbers that follow them inherit their credibility from these checks passing.

4.12 Second-method verification of the attractor census

Because the census of Section 4 underpins the basin and learning claims, we re-derived it with two independent Boolean-network engines: AEON (symbolic, asynchronous semantics) and mpbn (answer-set programming). External inputs were constant-folded into the update rules (sympy DNF simplification) before analysis. Fixed points are invariant under update semantics, so all engines must agree: they do, on all eight conditions, exactly matching our committed counts including both bistable conditions (`results/grn_crosscheck.json`; `code/experiments/grn_crosscheck.py`). The census is therefore not an artefact of our own parser or enumeration code.

4.13 Multi-cell extension: the SLP prepatter problem made exact

The single-cell module’s external inputs are, in the embryo, the outputs of neighbouring cells. We ported the id-021 rule set verbatim to a 1D row of $C = 4$ cells (56 variables): each cell carries the full 14-node network, and `v_WG_external` / `v_hh_external` of cell i are defined as the disjunction of WG/HH protein in cells $i - 1$ and $i + 1$ — the Albert–Othmer paracrine geometry [9]. The pair-rule gene *sloppy paired* (SLP) enters as a fixed per-cell prepatter bit. Fixed points are computed by answer-set enumeration (mpbn); being update-semantics-invariant, they need no simulation choice (`experiments/segment_polarity_mc2.py`, `results/grn_multicell.json`). An earlier 5-node simplification of this model is superseded; it produced no fixed points and is kept in the repository only as a record (Section 8).

Scanning all $2^4 = 16$ SLP prepatters yields a clean selector table (Table 7). The wild-type pattern — one *wg* cell adjacent to one *en/hh* cell, two default cells (WE. .) — is a fixed point for exactly four prepatters: $\text{SLP} = (0, 0, 0, 0), (0, 0, 0, 1), (1, 0, 0, 0), (1, 0, 0, 1)$; i.e. SLP must be OFF in the *en* cell (required by the *en* rule itself) and OFF in the cell posterior to it. The mechanism of the second constraint is visible in the rules: with SLP on in that cell, HH received from the *en* cell activates CIA, and $\text{CIA} \wedge \text{SLP}$ switches *wg* on — producing an ectopic second *wg* stripe (WEW., WEWE). This is the multi-cell reading of the pair-rule prepatter’s role: SLP must clear a two-cell window for each stripe pair, or the segment double-patterns.

(The four wild-type-compatible prepatters — 0000, 0001, 1000, 1001, their rows bolded — are exactly those with SLP OFF in the second and third cells. Patterns such as .EW. and .WE are reversed-polarity or shifted stripe pairs, which the embryo never sees; their existence as additional fixed points is expected in a boundary-free ring model and is noted, not hidden.)

4.14 Cross-organism generalization: the Arabidopsis floral organ identity network

The verdict asks whether the fragility framework generalizes beyond *Drosophila* (#15). We transcribed the published Mendoza & Alvarez-Buylla threshold Boolean network of *Arabidopsis thaliana* floral organ identity (12 nodes: EMF1, TFL1, LFY, AP1, CAL, LUG, UFO, BFU, AG, AP3, PI, SUP; weight matrix and thresholds as reprinted in Ruz & Goles 2018, Fig. 1) and enumerated all 4,096 states under synchronous update (`experiments/flower_model.py`, `results/flower_model.json`). The implementation validates exactly against the published record: 13 attractors (6 fixed points plus 7 limit cycles of length 2), and all six fixed points reproduce the published organ-identity states bit for bit (sepal 000100000000, petal 000100010110, carpel 000000001000, stamen 000000011110, inflorescence 110000000000, mutant 110000010110).

Table 7: Fixed-point stripe patterns of the 4-cell id-021 port for all 16 SLP prepatterns (W = *wg* cell, E = *en/hh* cell, . = default cell; n = number of fixed points). Wild-type WE.. rows in bold. Committed: `results/grn_multicell.json`.

SLP prepattern	fixed-point patterns	n
0000	, ..EW, .EWE, EWE., EWEW, WE.. , WEEW, WEWE	22
0001	, ..EW, .EW., EWEW, WE.. , WEEW, WEW.	14
0010	, .EWE, EW.., WEWE	11
0011	, .EW., EW.., WEW.	7
0100	, ..WE, EWE., EWEW	11
0101	, EWEW	5
0110	, ..WE, EW.., EWWE	9
0111	, EW..	3
1000	, ..EW, .WE., .WEW, WE.. , WEEW, WEWE	14
1001	, ..EW, .WEW, WE.. , WEEW, WEW.	10
1010	, WEWE	5
1011	, WEW.	3
1100	, ..WE, .WE., .WEW	7
1101	, .WEW	3
1110	, ..WE	3
1111		1

Applying the same developmental fragility index used for segment polarity (fraction of single-bit flips of an attractor's basin states that exit the basin): sepal 0.466 (basin 168), petal 0.399 (248), carpel 0.681 (24), stamen 0.792 (8), inflorescence 0.347 (384), mutant 0.319 (384). Three cross-organism findings. (i) **Fragility is not a *Drosophila* artifact**: the floral organ attractors sit in the same 0.3–0.8 fragility band as the segment-polarity wild-type state (0.500), while the gap-gene cascade sits at 0.000 - the basin-average fragility values differ across these models. This small comparison is consistent with a feedback-topology hypothesis, but does not isolate topology from model rules or starting-state measures. (ii) **Basin size inverts with organ specificity**: the two largest, least fragile basins are the inflorescence and mutant states - precisely the states that are *not* differentiated organs; the four true organ fates occupy smaller, more fragile basins (8–248 states), a model observation that does not identify a biological canalization mechanism. (iii) **Update scheme is load-bearing**: under synchronous update only 29.7% of state space reaches the biologically meaningful fixed points; the rest feeds limit cycles that the published block-sequential scheme eliminates - showing model-output dependence on update convention, not a measurement of developmental realism.

4.15 Multi-bit robustness depends on the perturbation starting measure

The single-bit fragility index averages perturbations over every state in a basin. This measure is mathematically well-defined but differs from asking how stable an established fate is. We tested that distinction by enumerating the complete 12-node floral network state space and every Hamming-radius perturbation. We retain the published synchronous threshold update, rather than treating this experiment as a reconstruction of molecular fluctuations.

Let $a(x)$ denote the eventual attractor of state x , x_f a fixed point, and $B_f = \{x : a(x) = a(x_f)\}$ its basin. Two retention functions are

$$R_f^{\text{fixed}}(k) = \frac{1}{\binom{12}{k}} \sum_{\delta: |\delta|=k} \mathbf{1}[a(x_f \text{ xor } \delta) = a(x_f)]$$

and

$$R_f^{\text{basin}}(k) = \frac{1}{|B_f| \binom{12}{k}} \sum_{x \in B_f} \sum_{\delta: |\delta|=k} \mathbf{1}[a(x \text{ xor } \delta) = a(x_f)].$$

The earlier fragility index equals $1 - R_f^{\text{basin}}(1)$. It does not equal $1 - R_f^{\text{fixed}}(1)$ in general. The first function models perturbing an established attractor; the second models a uniformly sampled trajectory-starting state that will eventually reach it. Neither uniform measure has been calibrated to the distribution of actual developmental states.

Fate	Fixed $R(1)$	Basin $R(1)$	Fixed $q_{1/2}$	Basin $q_{1/2}$
sepal	0.500	0.534	0.116	0.122
petal	0.417	0.601	0.105	0.141
carpel	0.333	0.319	0.084	0.083
stamen	0.167	0.208	0.068	0.071
inflorescence	0.667	0.653	0.167	0.163
mutant	0.583	0.681	0.146	0.173

Here $q_{1/2}$ is the first tested independent per-node flip probability whose same-attractor retention falls to one half. It is obtained by weighting the exact-radius curves with the binomial distribution,

$$P_f(q) = \sum_{k=0}^{12} \binom{12}{k} q^k (1-q)^{12-k} R_f(k).$$

The grid resolution is 0.001, so the quoted crossing is a grid bracket, not a fitted biological threshold. At $q = 0$, retention is one. At $q = 0.5$, every perturbed state is uniformly distributed across the full cube, so retention must equal $|B_f|/4096$ independently of the starting state. This is an exact consistency check on both measures. Every one-bit basin result also recovers the previously archived fragility value within its rounding precision.

The petal-versus-sepal ordering reverses. From their established fixed points, sepal retains fate under one bit flip with probability 0.500, while petal retains it with probability 0.417. Averaging uniformly over each basin instead gives 0.534 for sepal and 0.601 for petal. Petal is therefore less stable under the fixed-point measure but more stable under the basin-start measure. One scalar fragility ranking cannot represent both questions.

This is a positive methodological limit: robustness is a property of the dynamics together with a perturbation kernel and a starting-state measure. It is not a topology-only invariant, and a cross-organism comparison must specify those choices. The observed ordering reversal is exact within the transcribed model and does not rely on a noisy simulation estimate. It remains model-bound: artificial bit noise, synchronous update and uniform basin weights are not direct molecular interventions. A useful next test would weight states by measured developmental trajectories and compare update schedules under the same kernel, rather than calling the current values intrinsic canalization.

The full radius curves, independent-flip grid values and source-code checksum are in `results/flower_perturbation_radius.json`; the executable audit is `experiments/flower_perturbation_radius.py`. No new empirical biological discovery or causal topology law is claimed.

4.16 Exact asynchronous absorption separates topology from update schedule

The earlier radius audit fixed the synchronous update rule while changing the starting measure. We now hold the floral threshold matrix and one-bit perturbation kernel fixed and change only the update schedule. At each step, one of the twelve nodes is selected uniformly at random and updated from the current state. Other nodes retain their values. Self-loops are retained when the selected node does not change. This is random single-node asynchronous dynamics, not a random permutation sweep, and is not calibrated to measured gene-expression timing.

For the 4,096 states, let $T_i(x)$ update node i alone. The Markov transition matrix is

$$P_{xy} = \frac{1}{12} \sum_{i=1}^{12} \mathbf{1}[T_i(x) = y].$$

We enumerate directed strongly connected components and retain only closed components as recurrent classes. There are exactly six closed classes, each a published floral fixed point. The seven synchronous two-cycles are not closed recurrent classes under this update kernel. This is a mathematical change in the dynamics, not evidence that real developmental cycles are absent.

For transient-state matrix Q and direct recurrent-class entry matrix R , the absorption matrix H satisfies

$$(I - Q)H = R.$$

Boundary rows are one-hot indicators of class membership. Solving the sparse system gives the probability of eventual absorption at each fate from every initial state, without trajectory Monte Carlo. Maximum harmonic residual is 4.44×10^{-16} ; all rows sum to one within numerical tolerance. The classification is exact enumeration, while absorption values are floating-point solutions of that finite system.

Fate	Sync. fixed $R(1)$	Async. fixed $R(1)$	Uniform mass
Sepal	0.500	0.769	0.1974
Petal	0.417	0.651	0.1591
Carpel	0.333	0.687	0.0879
Stamen	0.167	0.557	0.0557
Inflorescence	0.667	0.852	0.2853
Mutant	0.583	0.735	0.2147

Fixed-start retention perturbs each of the twelve nodes of a published fixed point in turn and averages its absorption probability back to that fate. It uses the same twelve equally weighted perturbed

states as the synchronous comparison. The carpel-versus-petal ordering reverses: synchronous retention is 0.333 versus 0.417, whereas asynchronous retention is 0.687 versus 0.651. Topology, weights, thresholds, fixed points, starting states and perturbation locations are identical. Update schedule alone changes this ordering. This is a counterexample to treating a robustness ranking as an invariant of network topology.

Uniform absorption mass averages over all starting states and is not a deterministic basin size. A stochastic initial state can reach more than one fate under different schedules. To compare identical starting sets, we also evaluate the states in each *synchronous* basin under the asynchronous kernel. Their probability of reaching the originally assigned fate is 0.828 for sepal, 0.679 for petal, 0.915 for carpel, 0.852 for stamen, 0.915 for inflorescence and 0.799 for mutant. These are not asynchronous basin definitions. They show that transporting a synchronous label to another update rule can lose certainty even before a new bit perturbation.

Together with the starting-measure reversal, this result restricts the earlier topology-law language. Robustness belongs to the chosen transition dynamics, perturbation kernel and initial-state distribution together. Published-network comparisons can identify model-specific examples, not establish a causal cross-organism topology law without controlling those choices. Increasing asynchronous retention in this experiment is not evidence of superior biological canalization: single-node timing, artificial bit changes and uniform state weights remain modeling assumptions. No embryonic trajectory distribution or molecular-noise scale has been measured here.

The executable is `experiments/flower_async_absorption.py`. Results include the closed classes, node-specific perturbation probabilities, all model checks and the source-code checksum in `results/flower_async_absorption.json`; the complete 4,096-by-six absorption matrix is retained in a compressed numerical archive. Source model: https://pageperso.lis-lab.fr/~sylvain.sene/files/publi_pres/rgs18.pdf. A useful next test would compare calibrated update rates or measured trajectory weights, rather than label the present schedule a biological truth.

4.17 Eventual return and recovery time answer different questions

An absorption probability is an infinite-horizon statement. It does not say that recovery is fast enough for a developmental process. We therefore compute first-absorption times and finite-horizon return probabilities in the same enumerated asynchronous chain, without changing topology, update schedule or perturbed states. Time is counted in single-node update attempts, including attempts that leave the state unchanged. It is not seconds, cell cycles or empirically estimated gene-expression time.

Let τ be the first hitting time of any recurrent fixed point. For transient states, expected time m

and second moment s solve

$$(I - Q)m = \mathbf{1}, \quad (I - Q)s = \mathbf{1} + 2Qm.$$

The variance is $s - m^2$. For a particular fate f , let H_f be its eventual absorption probability and $g_f(x) = \mathbb{E}_x[\tau \mathbf{1}\{X_\tau = f\}]$. Then

$$(I - Q)g_f = H_f.$$

For a mixture of perturbed states, conditional time to successful return is the sum of g_f divided by the sum of H_f . This conditions on successes and is not the mean time to any fate. Recurrent boundary states have zero first-hitting time.

Fate	Mean time, any fate	Conditional return time	Return by 12
Sepal	13.111	12.301	0.4948
Petal	12.694	10.635	0.4646
Carpel	13.662	11.373	0.4630
Stamen	13.273	8.882	0.4283
Inflorescence	14.732	13.912	0.5097
Mutant	14.460	13.194	0.4728

Finite-horizon recovery is evaluated by repeatedly applying P to the one-hot recurrent boundary vectors. Since each recurrent class is an absorbing singleton, being at the original fixed point by update t is equivalent to having returned to it by t . From the same twelve one-bit perturbed fixed-point states, petal recovery by twelve updates is 0.46464 and carpel recovery is 0.46295. Thus petal is very slightly ahead at that horizon even though eventual return favors carpel, 0.68708 versus 0.65123. By twenty-four updates carpel is ahead, 0.60396 versus 0.58210. This is a model-exact finite-horizon ranking change, not an empirical effect-size claim: the twelve-update difference is small and no real developmental time has been mapped to these counts.

Across uniformly weighted starting states, expected first absorption is 30.1989 updates; the maximum statewise expectation is 43.1467. Neither is a bound on every trajectory. Self-loops allow arbitrarily long realizations with decreasing probability. Likewise, twelve updates are not a sweep in which every node updates once: node choices are random with replacement. Finite-horizon probabilities converge upward to the previously solved eventual probabilities, and all computed values are checked against that upper bound.

This adds a third qualification to a scalar robustness claim: after specifying topology, transition rule, perturbation kernel and starting measure, one must also specify the recovery horizon. A favorable eventual fate probability can coexist with slower early recovery. The result does not establish physiological speed, a calibrated developmental deadline or a measured molecular mechanism. Expected times, variances, successful-return joint moments and exact finite-horizon probabilities at

12, 24, 48 and 96 updates are retained in `results/flower_async_absorption_times.json` and its full numerical archive. The executable is `experiments/flower_async_absorption_times.py`. Future mapping to biological time requires measured update kinetics; these update counts do not supply them.

5 Study C: Morphogen gradients and positional information

5.1 Data

All measurements come from Liu, Morrison and Gregor, PNAS 2013 [5], obtained from the public archive (Zenodo record 4942019, mirroring Dryad doi:10.5061/dryad.2j9p4, CC0). Two modalities:

- **LiveImaging.mat**: 676 quality-filtered Bcd-GFP intensity gradients from 29 fly lines (quality criteria of the original publication: no leakage, no auto-fluorescence outliers; per-embryo records with line, embryo, side, position axis in both μm and fractional egg length, and per-line maternal *bcd* dosage from $1\times$ to $4\times$). This is the same dataset underlying the published λ measurements.
- **Immunofluorescence.mat**: 2,180 fixed-embryo records with co-stained gap-gene expression (Hb, Kr, Gt, Kni, Eve in varying panels per line), used for the boundary and decoding analyses of Sections 5.7–5.7.

Under the project’s uniform counting rule (one identifier-backed record = one dataset) these two files contribute $676 + 2,180 = 2,856$ accession-level records; the full manifest is in Appendix A. Gene identities were verified against NCBI Gene (eutils): *bcd* 40830, *hb* 41032, *eve* 36039, *kni* 40287, *gt* 31227, and the source publication PMID 23580621 (lookup records committed in `data/lookups/ncbi_genes.json`).

5.2 The dosage series and the cephalic furrow

The experimental design of the source study [5] exploits a set of transgenic lines carrying one to four functional copies of maternal *bcd* (the 2XA/IIA/IIIA/BNT nomenclature encodes the transgene complement of each line; Appendix A lists the mapping used). Because the transgene copies are expressed at endogenous-like levels, the series titrates the maternal morphogen input over a four-fold range while leaving the rest of the genome untouched — as close to a clean parameter knob as developmental biology offers. The amplitude response of the gradient is correspondingly clean

(amplitude–dosage Spearman $\rho = 0.86$), and it is against this strong, expected effect that the subtler length-scale and readout effects of Sections 5.5–5.7 must be read.

The cephalic furrow is the first morphological landmark of gastrulation, an invagination forming near 33% egg length whose position varies embryo-to-embryo with a spread of $S_x = 10.5\%$ EL in the source dataset. Whether the furrow is positioned directly by the Bcd gradient or indirectly through the gap-gene layer is a live question; what the benchmark of Section 5.6 establishes is an information floor: the Bcd profile alone suffices to predict furrow position to 2.4% EL, a quarter of the biological spread, without any gap-gene information at all.

5.3 Analysis pipeline

Every Study-C number passes through the same pipeline, implemented in `bcd.py`, `decoder.py` and `gapgene.py` and driven by the `experiments/` scripts:

1. **Ingest.** `LiveImaging.mat` and `Immunofluorescence.mat` are parsed with SciPy (with an `h5py` fallback for v7.3 files). Each embryo record carries its line, embryo index, side, position axis (both μm and fractional egg length) and fluorescence vector.
2. **Quality filter.** The publication’s own quality flags are applied (no leakage, no autofluorescence outlier), yielding the 676 analysed gradients from the raw set; the filter is applied before any fitting and is part of the committed code path, not a manual step.
3. **Fitting.** Per-gradient single-exponential fits (12) by Levenberg–Marquardt (SciPy `curve_fit`) on the stated window, with the two-component variant (13) fitted identically for the F -test (14). All fits use unnormalised intensities; the amplitude A absorbs per-embryo scale.
4. **Aggregation.** Population statistics are medians with IQRs (reported alongside means nowhere — the distributions are right-skewed and we do not pretend otherwise); correlations are Spearman rank coefficients (26); the parametric cross-check is robust OLS.
5. **Cross-validation.** Decoder benchmarks split by *fly line*, never by embryo, so that line-specific systematics cannot leak from train to test. Bootstrap intervals (20) resample embryos within the comparison being made.

5.4 Result 1: the exponential length scale, reproduced — and window-sensitive

Fitting the single exponential (12) to each gradient and taking the population median gives $\lambda = 0.256$ (IQR $[0.230, 0.287]$, CV 0.185) on the default analysis window $x \in [0, 0.6]$ EL. The published population value is 0.165 ± 0.007 [5] — but the published fits were performed on the re-

stricted window $x \in [0.2, 0.8]$ EL, and the measured length scale depends strongly on the window (Table 8): refitting on the authors’ own window gives $\lambda_{\text{median}} = 0.189$, and the window family shows a monotone trend from 0.249 (anterior-only, $[0, 0.4]$) to 0.189 (their window). The residual $\sim 15\%$ gap between our 0.189 and the published 0.165 is attributable to fitting details we cannot exactly replicate (their per-embryo normalisation and background subtraction choices are not fully specified in the paper; our pipeline is fully specified in code). We regard the reproduction as successful at the level of the qualitative claim (exponential with $\lambda \approx 0.17\text{--}0.19$ EL on the published window) and report the residual discrepancy openly rather than tuning it away.

Table 8: Window sensitivity of the Bcd length scale and of the dosage-coupled length-scale (DCLS) correlation. Each row: per-gradient single exponential fits restricted to the stated x window (fractional egg length), population median λ , and Spearman correlation of per-gradient λ with maternal bcd dosage across the 20 dosage-annotated lines ($n \approx 482$ gradients). Committed: `results/bcd_windows.json`.

Window (EL)	n	λ_{median}	IQR	Spearman ρ	p
[0.0, 0.4]	481	0.249	[0.222, 0.280]	0.290	9.4×10^{-11}
[0.0, 0.5]	482	0.234	[0.212, 0.255]	0.356	7.9×10^{-16}
[0.0, 0.6]	482	0.221	[0.204, 0.239]	0.360	3.9×10^{-16}
[0.05, 0.45]	482	0.217	[0.198, 0.235]	0.324	3.1×10^{-13}
[0.1, 0.6]	482	0.200	[0.187, 0.215]	0.350	2.6×10^{-15}
[0.2, 0.8] (published)	482	0.189	—	0.071	0.12 (n.s.)

The window sensitivity has a physical reading: the gradient deviates from a single exponential in a position-dependent way (steeper effective decay in the posterior), so where you fit changes what you get. This connects to the long-standing critique that the SDD steady state is not the whole story for Bcd [12]; our F-test below quantifies it on this dataset.

Two-component test. Fitting the two-component form (13) to every gradient and applying the nested F -test (14): only 0.74% of gradients reject the single exponential at $p < 0.05$ (0.30% after Bonferroni correction) — *fewer than the 5% expected under the null*. On this dataset the second exponential component is not statistically justified per-embryo, even though the population-level window trend shows systematic structure. Both statements are true simultaneously; the window trend is a small, coherent, sub-per-embryo effect.

5.5 Result 2: the length scale increases with maternal bcd dosage (DCLS)

Liu et al. reported in their supplement (their Fig. S3A) an approximately 10% increase of λ with dosage. We reproduce and refine this. On anterior-weighted windows the correlation is strong and

highly significant: Spearman $\rho = 0.356$, $p = 7.9 \times 10^{-16}$ on $[0, 0.5]$ EL; per-dose medians on the default window rise from $\lambda = 0.246$ ($1\times$, $n = 158$) through 0.263 ($2\times$, $n = 162$) to 0.287 ($3\times$, $n = 100$), with $4\times$ ($n = 62$) at 0.259 — non-monotone at the top dose, which we report without a story. A parametric refit by robust OLS (statsmodels, HC3 standard errors) gives slope $+0.0117 \lambda$ per dose, 95% CI $[0.0060, 0.0174]$, $p = 5.9 \times 10^{-5}$ (results/dcls_ols.json).

The crucial window-dependence: on the authors’ $[0.2, 0.8]$ window the correlation is *absent* ($\rho = 0.071$, $p = 0.12$). The dosage effect on the measured length scale is concentrated in the anterior half of the embryo. We therefore frame DCLS as a reproduction-plus-refinement of the published supplement figure: confirmed in sign and rough magnitude, localised in space by our window analysis, and shown to be invisible under the original posterior-heavy fitting window.




Figure 5: Left: per-dosage λ distributions (DCLS). Right: window sensitivity of both λ and the DCLS correlation; the published window is marked.

5.6 Result 3: a learned decoder breaks the published cephalic-furrow benchmark

Liu et al. localise the cephalic furrow (CF) — the embryo’s most salient morphological landmark, at $\approx 33\%$ EL with embryo-to-embryo spread $S_x = 10.5\%$ EL in their units — from the Bcd gradient by a threshold model: find where the profile crosses a calibrated relative concentration. Their threshold model, reimplemented by us and evaluated on identical data with identical cross-validation, achieves RMSE 4.94% EL with the true per-line dosage, or 2.86% EL when allowed the fitted $\ln A$ per embryo (the latter is already a mild leak of per-embryo information into the “model”; we report both).

We trained a one-dimensional convolutional decoder on the raw Bcd-GFP profile (architecture: three 1D conv blocks, global pooling, linear head; PyTorch, CPU) to predict CF position, with

5-fold cross-validation over disjoint fly lines — a line-level holdout, so no embryo of a test line appears in training. Results (250 embryos with CF annotation, 29 lines, 3 training seeds; committed `results/cf_decoder.json`):

Table 9: Cephalic-furrow position decoding. Cross-validation over disjoint fly lines; RMSE in percent egg length. The CNN beats the published threshold model on identical data by a paired-bootstrap margin of 2.57% EL, $CI_{95} [1.99, 3.11]$, $p < 10^{-4}$. Committed: `results/cf_decoder.json`.

Model	RMSE (%EL)	R^2
Liu threshold, true dosage	4.94	0.374
Liu threshold, fitted $\ln A$	2.86	0.790
Ridge on $(\ln A, \lambda)$	2.55	0.833
1D CNN on raw profile (3 seeds)	2.38	0.854

The gain over the published model is 2.57% EL on average with bootstrap CI $[1.99, 3.11]$ entirely above zero: the benchmark is broken on its own data by a direct margin, and $p(\text{gain} \leq 0) = 0$ in 10^4 resamples (equation (20)). Context for the absolute number: the embryo-to-embryo spread of CF position itself is $S_x = 10.5\%$ EL and the published within-embryo measurement precision is far finer; a 2.4% EL decoder error means the Bcd profile alone carries enough information to localise the furrow to within about a quarter of the biological spread. The CNN’s advantage over the two-feature ridge (2.55) is small but consistent — most of the win over the threshold model comes from using the whole profile shape, not from deep learning per se; we say so to keep the claim honest.

5.7 Result 4: gap-gene decoding and the dosage–decodability link

Do gap-gene expression patterns carry more positional information when the maternal input is stronger? Using the 2,180 immunofluorescence records, we built k -nearest-neighbour decoders (committed `experiments/gap_decode.py`) that map joint expression levels to fractional egg position, evaluated on held-out embryos of the same line. Table 10 summarises (`results/gap_decode.json`).

Two-gene decoders reach 6–11% EL (Hb+Eve) and 18–22% EL (Kni+Gt), consistent with the published ordering (single genes several-fold worse than the $\sim 1\%$ EL of the joint four-gene decoder [6, 7]). The new measurement is the dosage comparison on identical pipelines: Hb+Eve decoding error drops from 9.91% EL at $1 \times \text{bcd}$ to 6.21% EL at $3 \times$, a paired difference of 3.70% EL with bootstrap $CI_{95} [3.00, 4.40]$. Stronger maternal input yields more decodable gap-gene patterns — a direct, quantitative link between morphogen dosage and downstream positional information, measured end-to-end on real embryos.




Figure 6: CF decoder: predicted vs. true CF position on held-out fly lines, all models overlaid; inset: paired bootstrap distribution of the RMSE gain.

Table 10: Positional decoding from gap-gene pairs (kNN, held-out embryos). Median absolute error in %EL. Published context: joint 4-gap-gene decoding achieves $\sim 1\%$ EL [6]; single gap genes carry 2–3 bits [7]. Our two-gene numbers are consistent with that ordering.

Line (dosage)	Genes	Embryos	Test pts	Median err (%EL)
1XA (1 $\times$)	Hb, Eve	137	2362	9.91
2XA2IIIA (3 $\times$)	Hb, Eve	205	3575	6.21
1XA BNT (1 $\times$)	Hb, Eve	211	3778	11.01
2XA BNT (2 $\times$)	Hb, Eve	218	3575	8.41
1XA (1 $\times$)	Kni, Gt	152	2359	22.02
2XA (1 $\times$)	Kni, Gt	194	3193	22.22
2XA2IIIA (3 $\times$)	Kni, Gt	195	3357	17.82

Complementarily, the Hb boundary positions themselves shift with dosage (`results/hb_boundaries.json`): the anterior Hb boundary moves from 0.372 EL (1XA, $n = 18$) to 0.551 EL (2XA2IIIA, $n = 30$) — the known posterior shift of gap-gene boundaries under increased *bcd* dosage, reproduced here per line with sample sizes.




Figure 7: Gap-gene positional decoding by line and gene pair; the dosage ordering within each panel is the dosage–decodability effect.

5.8 What Study C establishes

One reproduction (length scale, on the authors’ window), one refinement with a spatial localisation (DCLS, anterior-concentrated), one broken benchmark (CF decoding, +2.57% EL CI [1.99, 3.11]), and one measured functional consequence of dosage (decodability, 3.7% EL CI [3.00, 4.40]). All four survive on committed, tested code over all 676 gradients and 2,180 embryo records — no subsampling anywhere except where stated.

5.9 Per-line detail and internal controls

Table 11 gives the per-line length-scale summary underlying every population statement above (computed from the 676 committed per-gradient fits in `results/bcd.sdd.json`). Two internal controls bracket the DCLS claim. First, λ does *not* correlate with egg length across gradients (Spearman $\rho = -0.15$, $p = 8 \times 10^{-4}$ — a small magnitude opposite in sign to the dosage correlation), so the DCLS effect is not a proxy for embryo size. Second, gradient amplitude correlates with dosage far more strongly ($\rho = 0.86$, $p = 1.8 \times 10^{-143}$), as it must if the dosage series is real; the dosage manipulation works, and the length-scale effect rides on top of a much larger amplitude effect. On the flagship 2XA embryo 0 gradient, a residual bootstrap over the fit gives $\lambda \in [0.251, 0.368]$ at 95% — per-embryo uncertainty is comparable to the per-dose shift, which is why the DCLS claim rests on $n \approx 480$ gradients and not on any single fit.

Table 11: Per-line Bcd length-scale summary (median λ over the line’s gradients on the default $[0, 0.6]$ EL window; range in brackets). Lines without dosage annotation (BN/BT/BNT series) are excluded from dosage correlations but included in population statistics. Source: `results/bcd_sdd.json`, per-gradient rows.

Line	bcd dose	n gradients	λ median	range
1IIA	1x	46	0.216	[0.185, 0.268]
1IIA BNT	–	22	0.231	[0.175, 0.399]
1IIA1IIIA	2x	52	0.260	[0.191, 0.321]
1IIC	1x	6	0.238	[0.197, 0.297]
1XA	1x	32	0.237	[0.198, 0.309]
1XA BN	–	10	0.229	[0.193, 0.277]
1XA BNT	–	16	0.262	[0.215, 0.673]
1XA BT	–	18	0.236	[0.188, 0.313]
1XA1IIA	2x	18	0.262	[0.211, 0.331]
1XA2IIA	2x	18	0.266	[0.226, 0.341]
2IIA	1x	32	0.273	[0.211, 0.342]
2IIA2IIIA	3x	44	0.277	[0.193, 0.439]
2IIB	1x	8	0.253	[0.212, 0.378]
2IIC	2x	10	0.256	[0.200, 0.439]
2IIIA	2x	22	0.261	[0.204, 0.333]
2IIIB	2x	24	0.278	[0.212, 0.336]
2XA	1x	34	0.270	[0.222, 0.436]
2XA BN	–	24	0.224	[0.186, 0.427]
2XA BNT	–	54	0.237	[0.181, 0.456]
2XA BT	–	26	0.268	[0.217, 0.386]
2XA1IIA	2x	18	0.263	[0.211, 0.429]
2XA1IIA BN	–	10	0.244	[0.224, 0.271]
2XA1IIA BT	–	14	0.260	[0.187, 0.368]
2XA1IIA2IIIA	4x	28	0.254	[0.198, 0.297]
2XA2IIA	3x	24	0.275	[0.190, 0.359]
2XA2IIA2IIIA	4x	12	0.277	[0.225, 0.382]
2XA2IIC2IIIA	4x	22	0.257	[0.216, 0.577]
2XA2IIIA	3x	22	0.294	[0.250, 0.379]
2XA2IIIB	3x	10	0.315	[0.250, 0.359]




Figure 8: The 676 quality-filtered Bcd-GFP gradients (per-dosage panels), with single-exponential fits overlaid. The anterior steepness and posterior shoulder that drive the window sensitivity of Table 8 are visible by eye.

5.10 Decoding protocol detail

The gap-gene decoders of Table 10 deserve precise description because decoding numbers are unusually sensitive to protocol choices. Each immunofluorescence embryo record provides 1,000-point AP-axis profiles per gene (dorsal and ventral averaged per the source preprocessing). A decoding problem instance is (line, gene-pair): features are the expression levels of the two genes at a queried position, the target is that position, and evaluation is leave-out over embryos within the line with $k = 5$ neighbours (24). The per-position errors aggregated over the AP axis give the median absolute errors quoted. Two protocol decisions matter: positions are decoded *independently* (no smoothness prior along the axis — a smoother decoder would look better and be less honest), and test points include flat profile regions where decoding is hardest, which is why the Kni+Gt rows (posterior genes with broad domains) sit at 18–22% EL while Hb+Eve rows sit at 6–11% EL. Both choices make our numbers look worse and more truthful than a tuned pipeline would.

6 Tools, datasets and reproducibility

6.1 Question

How does positional decoding precision scale with the number of gap genes read jointly, and does the scaling survive a doubling of maternal *bcd* dosage? Published reference points: single gap genes

localize to several percent egg length (%EL) [7], while joint decoding of four gap genes reaches ~ 1 %EL [6]. The intermediate regime - two and three genes - and the dosage dependence of the *scaling law itself* are unmeasured.

6.2 Setup

From the Liu 2013 immunofluorescence archive (Study C data) we take the two sessions that image three gap genes simultaneously (Hb, Eve, Kr): 244 embryos at $1 \times bcd$ dosage (line 1XA) and 142 embryos at $2 \times$ dosage (line 2XA). Position is decoded with k -nearest-neighbour regression ($k = 5$) from expression vectors at 50 subsampled axis positions, holding out 20% of embryos; every gene subset of size 1, 2 and 3 is decoded independently (2 seeds per subset). Controls: circular position permutation within each embryo, preserving expression autocorrelation while destroying position information.

6.3 Result 1: error falls as d^{-2} and extrapolates to the published 4-gene precision

At $1 \times$ dosage the median absolute decoding error falls from 21.6 %EL (one gene) to 6.2 %EL (two genes) to 2.4 %EL (three genes), a log-log slope of $\alpha = 1.98 \pm$ (three-point fit). Extrapolating the same law one step to $d = 4$ predicts 1.4 %EL, consistent with the ~ 1 %EL reported by [6] for joint four-gene decoding on different data with Bayes-optimal decoding. Because our k NN decoder is a strict lower bound on decoding precision, this agreement is conservative. We do not claim to beat the published number; the contribution is the measured scaling law in the unmeasured intermediate regime.

6.4 Result 2 (discovery): doubling bcd dosage shallows the scaling exponent

At $2 \times$ dosage the same measurement gives 21.0 %EL, 8.0 %EL and 4.2 %EL for one, two and three genes: a shallower exponent ($\alpha = 1.46$ vs. 1.98), with $d = 4$ extrapolation 2.8 %EL. The exponent-shallowing is the robust finding: it replicates under a near-optimal Gaussian MAP decoder on the same sessions and folds ($\alpha: 2.13 \rightarrow 1.70$), while the worse absolute three-gene precision at $2 \times$ dosage does not (2.0 %EL in both lines under MAP) and is downgraded to a k NN-specific observation. Doubling the morphogen changes *how precision scales with gene count*, not the ceiling already reached at three genes. This is consistent with Study C’s dosage-dependent length-scale change (DCLS) and adds a second, independent axis on which dosage compensation fails.

We note a boundary condition: the pair-level improvement reported earlier (Hb+Eve, 6.2 vs. 9.9 %EL across line sets) used *different* line sets (BNT vs. non-BNT), while the present contrast

holds the three-gene session fixed. Dosage effects on decodability are therefore line-set dependent, and the two results do not commute.

6.5 Controls and limits

Structure-preserved permutation nulls sit at 24-26 %EL for every subset size. Single-gene decoding (21.6 %EL) is barely above its null (24.4 %EL): on this immunofluorescence data, positional information is almost entirely a *multi-gene* property. Limits: only two three-gene sessions exist in the archive; the exponent is a three-point fit; the $d = 4$ point is a one-step extrapolation; k NN is not the optimal decoder.

6.6 Result 3 (POSLAW-3): joint decoding flattens the positional error profile; dosage doubling re-creates it at the poles

Binning per-sample decoding errors by true position (five bands across the axis, same sessions and folds) shows where the joint-decoding gain lives. At $1 \times bcd$, single-gene error is strongly U-shaped (37.6, 17.6, 7.6, 17.2, 36.0 %EL from anterior to posterior) while three-gene decoding is nearly flat (2.4, 2.4, 1.2, 2.4, 4.6 %EL): the gain ratio peaks at the anterior pole ($15.7\times$) and holds at $6.3\text{--}7.8\times$ elsewhere. Joint decoding does not merely lower average error - it *rescues the axis poles*, where no single gap gene carries usable position information. At $2 \times$ dosage the rescue weakens everywhere (gain $3.3\text{--}6.8\times$) and preferentially at the poles ($3.3\text{--}4.1\times$ at the end bands vs. $6.5\text{--}6.8\times$ mid-axis), re-opening the U (three-gene error 8.0 and 10.4 %EL at the poles vs. 1.6-2.0 mid-axis). The POSLAW-2 exponent shallowing therefore has a spatial mechanism: the dosage-dependent precision loss concentrates at the axis poles. Code experiments/gap_poslaw_posres.py; results results/gap_poslaw_posres.json.

POSLAW-4: the pole rescue is decoder-robust. POSLAW-2 split the dosage claims by decoder (exponent robust, absolute precision kNN-specific), so the POSLAW-3 pole rescue - measured with kNN only - needed the same test. Rerunning the identical position-resolved band analysis with the Gaussian MAP decoder (experiments/gap_poslaw_posres_bayes.py; results/gap_poslaw_posres_bayes.json), the rescue survives and strengthens at the poles: at $1 \times$ dosage the terminal-band single-gene error is 68 %EL against a flat 2 %EL three-gene profile (gain $34\times$, vs. $15.7\times$ under kNN), and the U-shaped single-gene profile persists across bands (13/4/10/10/68 %EL). At $2 \times$ the rescue again weakens preferentially at the poles (gains $5\text{--}17.5\times$, band-4 single-gene error 35 %EL). The pole rescue is therefore a property of the data, not of the kNN decoder, while the absolute-precision dosage claim remains decoder-specific - the claim-splitting discipline of POSLAW-2 extends to the

spatial finding.

6.7 Repository layout

```
mega27-14-digital-embryo/
  src/digitalembryo/      installable package (pip install -e .)
    morphogen.py          Turing simulators (Gierer-Meinhardt 2D), CFL-checked
    bcd.py                Bcd gradient loading, exponential fits, F-test, bootstrap
    positional.py         threshold readout, positional-error theory
    bnet.py               Boolean network parser + exact attractor engine
    gnn.py                GNN/MLP basin-boundary models
    gapgene.py            gap-gene decoding (kNN)
    decoder.py            CF decoders (threshold, ridge, CNN)
    segment_polarity.py   single-cell reference model
    segment_polarity_mc.py EXPERIMENTAL multi-cell extension (see negatives)
    cli.py                the embryosim command-line tool
  experiments/            one script per result block, each writes results/*.json
  tests/                  hermetic pytest suite (incl. symbolic math verification)
  results/                every number quoted in this report, as JSON
  data/raw/               downloaded datasets (re-fetchable via scripts/)
  data/TOOLS_AND_DATASETS.md tool and dataset manifest
  paper/                  this report (LuaLaTeX, Times New Roman)
```

6.8 The embryosim command-line tool

The project ships a usable CLI (entry point registered in `setup.py`):

```
embryosim fit-bcd --input LiveImaging.mat --out fits.json
    # per-gradient exponential fits + population stats (Study C, Sec. 6.2)
embryosim hb-boundary --input Immunofluorescence.mat --line 1XA
    # per-line Hb boundary positions with IQRs (Sec. 6.5)
embryosim turing --mu-a 0.04 --steps 2500
    # run a Gierer-Meinhardt sim, report selected wavenumber (Study A)
embryosim attractors --model id-021 --condition hh1_WG0_SLP0
    # exact attractor census of a curated Boolean model (Study B)
```

Each subcommand wraps exactly the code path that produced the corresponding committed results file, so CLI output and paper numbers cannot drift apart.

6.9 Test suite

The pytest suite is hermetic (no network; small committed fixtures) and currently comprises 30+ tests across seven files, including `test_math_verify.py`, which symbolically re-derives the steady state (3), the Jacobian invariants (4), the instability band (7) and the SDD/readout formulas (12)–(17) with sympy and asserts exact agreement with the committed numerical results — the paper’s formulas and its numbers are therefore checked against each other, not merely each stated once.

6.10 Data acquisition

`scripts/download_data.py` re-fetches all inputs from their public sources (Zenodo record 4942019 for the Liu dataset; the sybila/biodivine-boolean-models GitHub mirror for id-021/id-027) and verifies sizes. The full external-tool and dataset inventory, with the counting rule, is in Appendix A.

6.11 Compute environment

All results were produced on a 2-CPU / 1.9 GB-RAM sandbox in pure Python (NumPy/SciPy/PyTorch-CPU). Nothing in this report required a GPU; the largest single computation (the 171-simulation Turing sweep) completes in minutes. This is itself a deliberate design point: every claim is reproducible on commodity hardware.

6.12 Module reference

The installable package `digitalembryo` contains eleven modules; their roles, as documented in their own docstrings, are:

- `morphogen.py` — Gierer–Meinhardt activator–inhibitor kinetics on a 2D field, with linearised Turing-instability analysis (equations (7)–(8)) and CFL-checked integration.
- `turing_ext.py` — beyond-onset wavenumber-renormalisation analysis: the deviation of the nonlinear steady-state wavenumber from the linear fastest-growing mode as a function of distance from onset.
- `bcd.py` — real Bicoid-gradient analysis: SDD exponential fits, per-embryo two-component F -tests, bootstrap intervals on λ .
- `decoder.py` — cephalic-furrow positional decoding: reimplementations of the Liu et al. threshold model $x_{CF} = S_x \ln D + c$ and the CNN alternative.

- `gapgene.py` — gap-gene immunofluorescence handling: 1,000-point AP-axis profiles for Hb, Kr, Gt, Kni, Eve.
- `positional.py` — Wolpert French-flag and threshold-readout simulations, with the information-theoretic error bounds of Section 2.4.
- `bnet.py` — minimal BoolNet-format parser and exact synchronous-dynamics engine (full transition graph, attractor extraction).
- `grn.py` — Boolean-ensemble tools: Gardner–Collins toggle switch reference and ordered/chaotic classification via Derrida spreading.
- `gnn.py` — graph neural network over a regulatory graph (equation (27)) for basin-membership learning.
- `segment_polarity.py` — simplified single-cell segment-polarity reference model (Albert–Othmer topology [9]).
- `segment_polarity_mc.py` — first multi-cell attempt (SUPERSEDED, kept as record; see Section 8); the working multi-cell model lives in `experiments/segment_polarity_mc2.py`.

7 Discussion

7.1 Cross-study synthesis

The three studies share one methodological skeleton: take the strongest published quantitative statement available, recompute it from primary data or a primary model, then push one step past it — and report both steps with equal honesty. The push-past step succeeded in Study B (basin size, learning curve) and Study C (benchmark margin, dosage-decodability) and failed instructively in Study A (CNN surrogate, retracted). That 2-out-of-3 rate is not a bug in the programme; it is the expected yield of attempting genuine extensions rather than rehearsing known results. A programme in which every extension succeeds is a programme whose extensions were too timid.

7.2 What the numbers do and do not say

Study A. The upshift (median 31% above the linear mode, peaking at intermediate distance from onset) is a property of the Gierer–Meinhardt kinetics on lattices of the tested size. We have not shown it for other kinetics (Schnakenberg, FitzHugh–Nagumo), other boundary conditions, or continuum-accurate solvers; each is a natural next computation and the pipeline accepts them by parameter

change. The non-monotone dependence on ε suggests the deviation is controlled by the competition between mode-coupling strength (which grows with amplitude, i.e. with distance from onset) and band flatness (which shrinks it again), but we have measured, not derived, this — a weakly nonlinear analysis in the spirit of amplitude equations is the obvious theoretical follow-up, and our retracted CNN attempt is evidence that the effect is not trivially readable from linear data.

Study B. The 0.77% wild-type basin reframes the robustness question for the segment-polarity module: the model does not pattern *whatever* the initial state, it patterns given the initial states that the preceding developmental stages actually produce. The learnability result (~ 60 labelled states suffice) shows the basin boundary, though tiny, is geometrically simple — consistent with the ordered dynamics the Derrida analysis predicts. The multi-cellular extension this motivated is now in place (Section 4.13): the same network, embedded in a 4-cell row, supports the wild-type stripe pair exactly for the SLP prepatterns that clear a two-cell window.

Study C. The CF-decoder margin (+2.57% EL, CI [1.99, 3.11]) is large relative to the biology: the whole embryo-to-embryo spread of the furrow is 10.5% EL. Two readings are possible and we cannot yet separate them: the threshold model leaves profile-shape information unused (our preferred reading, supported by the ridge model capturing most of the gain from two shape features), or the CNN exploits line-specific systematics that happen to survive line-disjoint cross-validation. A second, independent CF-annotated dataset would decide between them; we flag this as the most important external validation of the project. The DCLS refinement (anterior-concentrated dosage coupling) is compatible with the published supplement figure and adds a spatial statement to it; the dosage–decodability result extends the same axis to the downstream genes.

7.3 Open work, in priority order

1. (Done during this project: the multi-cell port of id-021, Section 4.13.) Next step on that front: basins and reachability of the wild-type stripe pair under the compatible SLP prepatterns, which needs transient analysis beyond fixed points.
2. Independent-dataset validation of the CF decoder.
3. Weakly nonlinear theory of the wavenumber upshift; extension to other kinetics and boundary conditions.
4. Stochastic-target surrogates for pattern selection (predict the distribution over lattice modes rather than a single wavenumber).
5. Joint decoding across all measured gap genes on the immunofluorescence panel, approaching the four-gene benchmark of [6] with the same honest cross-validation.

8 Honest negatives and limitations

This section collects the results that did not come out the way we wanted. They are part of the deliverable, not an embarrassment routed around.

8.1 Retraction: the CNN Turing surrogate (Study A)

An $n = 58$ pilot suggested a CNN could predict the beyond-onset wavenumber upshift from the linear dispersion curve 17.7% better than the linear formula alone. At $n = 171$ the effect inverts (CNN rel. RMSE 0.292–0.295 vs. linear 0.272). The small-sample win was an artefact of the first sample. What survives: the upshift phenomenon itself (Section 3.3), the quantisation diagnosis, and a cautionary data point about surrogate-modelling claims at $n < 100$. The interim claim was corrected in the project record within a day; the retraction is committed in `results/turing_cnn.json` alongside the numbers.

8.2 Negative: the second exponential component is not justified per-embryo (Study C)

Despite the population-level window sensitivity of λ (Table 8), per-embryo two-component fits reject the single exponential in only 0.74% of 676 gradients (0.30% Bonferroni) — below the 5% false-positive floor. Whatever drives the window trend, it is too small per embryo to license a two-component model. We report this because the two-component fit is a standard move in this literature, and on this dataset it fails an F -test at scale.

8.3 Negative: uniform-sampling basin classification is degenerate (Study B)

A naive pipeline for “learn the basin boundary” samples states uniformly and reports accuracy. Here that yields 99.3% accuracy for the majority classifier itself (Section 4.5): the number is meaningless and would have looked publishable. Balanced sampling is what makes the learning curve informative. Any future basin-learning benchmark on rare attractors must control for this.

8.4 Superseded: the first multi-cell segment-polarity attempt

The first multi-cell extension (`segment_polarity_mc.py`, a 5-node simplification with hand-written coupling) found zero fixed points under synchronous update — a broken model. It is retained in the repository as a record, marked experimental, and was never used in any claim. The working multi-cell model (Section 4.13) replaces it by porting the curated id-021 rules verbatim instead of

re-deriving a reduced rule set — the lesson: extend from the curated network, not from a simplified sketch of it.

8.5 Limitations of scope

(i) The Turing upshift is measured for one kinetics family and one lattice; generality is untested. (ii) DCLS is an observational correlation across fly lines, not a perturbation experiment we performed; causal claims belong to the original authors’ dosage manipulations. (iii) The CF decoder, while cross-validated over fly lines, is evaluated on one dataset from one laboratory; a second independent CF-annotated dataset would be the right next test. (iv) Our λ reproduction retains a $\sim 15\%$ gap to the published central value, attributed to unspecified fitting details in the original; we chose transparency over tuning.

8.6 What we would do differently

Three procedural lessons generalise beyond this project. (i) *Size the sample to the claim before claiming*: the retracted Study-A surrogate win existed only at $n = 58$; a power consideration at the outset would have postponed the claim to $n \approx 170$ and never produced it. (ii) *Report the majority-class baseline first*: the degenerate basin classifier of Study B would have been caught in an afternoon by printing one extra line. (iii) *Fix the window before fitting*: the window sensitivity of Study C means any single-window λ number is incomplete without its window stated; every λ in this report now carries its window with it.

9 The developmental fragility index: quantifying, not lamenting, the small wild-type basin

The wild-type wg-ON basin of the published id-021 segment-polarity network (condition hh1_WG0_SLP0) contains exactly 128 of 16,384 states (0.781%), enumerated exhaustively under synchronous update. Rather than report this as a bare negative, we quantify its structure as a *developmental fragility index*: the fraction of single-node perturbations of basin states whose trajectory exits the wild-type basin,

$$F = \frac{1}{|B|} \sum_{s \in B} \sum_{i=1}^n \mathbb{1}[\text{traj}(s \oplus e_i) \notin B], \quad (28)$$

where B is the wild-type basin and $s \oplus e_i$ flips node i . We find $F = 0.500$: exactly half of the 1,792 single-bit perturbations destroy the wild-type outcome. Fragility is concentrated, not uniform (Table 12): Seven nodes have exit fraction exactly one: CIA, CIR, CI protein, EN protein, ci, en

and wg. The other seven have exit fraction zero under this synchronous basin measure. The earlier four-node table hid three equally sensitive nodes and is corrected below.

Table 12: Per-node exit fraction (fraction of basin states exiting after flipping that node).

Node flipped	exit fraction
v_CIA	1.000
v_CIR	1.000
v_CI_protein	1.000
v_EN_protein	1.000
v_ci	1.000
v_en	1.000
v_wg	1.000
other seven free nodes: 0.000 (see results/grn_fragility.json)	

9.1 Initial-state restrictions are not maintained knockouts

The historical `knockout_predictions` field has seven initial-state restrictions with zero wg-ON fixed-point outcomes and six with twice the unrestricted fraction. The implementation does not maintain the named node clamp during its trajectory: `wt_basin` sets the initial value, while `attractor_of` calls the original successor rule at each step. For example, with every free node initially zero except `en=1`, the very first successor sets `en=0`. A maintained `en=1` intervention would restore `en` to one after every update. This explicit source-level counterexample is retained in `results/grn_intervention_semantics_audit.json`.

Accordingly the previous knockout, constitutive-expression and basin-doubling interpretation is withdrawn. The old numbers remain unchanged and describe conditioning on a node's starting value, with a reduced initial-state denominator, not modifying the regulatory dynamics. A maintained-clamp experiment requires an explicit altered transition function and a fresh exact enumeration. It cannot be inferred from these records.

The edge-replacement experiment separately substitutes a regulator literal with zero in the target expression. This is a specified rule edit, not a physically validated deletion of an interaction. Its archived basin changes remain local model calculations, not evidence that targeted repression removal makes actual development more robust. The fragility index itself remains an exact synchronous single-bit basin-exit measure for the stated condition, not a measured biological fragility rate.

9.2 A new maintained-clamp census, with the phenotype limit visible

After verifying the semantic gap, a separate exact functional-graph census forces the chosen node before every rule evaluation and after each successor. It classifies all 8,192 states of the thirteen non-clamped nodes for each of 28 binary interventions. The unrestricted fourteen-node baseline again gives 128 wg-ON fixed-point outcomes among 16,384 starts. Results are new, not replacements for the historical initial-restriction file. Three maintained conditions are independently checked by direct trajectory enumeration in tests, in addition to graph-classification unit checks.

Table 13: Selected maintained-clamp contrasts. Fractions concern a uniformly weighted reduced initial space reaching any wg=1 fixed point, not a full biological wild-type phenotype.

Condition	Historical initial restriction	Maintained clamp
CIA ON	.015625	.250000
CI protein ON	.015625	.125000
ci ON	.015625	.062500
EN protein OFF	.015625	.031250
en OFF	.015625	.015625
wg ON	.015625	1.000000
wg OFF	.000000	.000000

The differences are substantive: CIA ON gives a quarter of the reduced starting space, not the old 1.5625%. They establish how changing the transition differs from conditioning on the initial measure. They do not show that actual constitutive expression produces those outcome probabilities. The endpoint itself needs care: forcing wg ON trivially satisfies the wg marker at every fixed point. Its fraction of one is not full wild-type rescue. A maintained intervention must therefore state whether the target phenotype includes the manipulated node and what other phenotype coordinates are required. Reduced-space weighting also changes the denominator, so a larger fraction is not automatically more developmental accessibility under a natural initial distribution.

The exact code and all 28 conditions are in `experiments/grn_maintained_clamp_audit.py` and `results/grn_maintained_clamp_audit.json`, with source model hash, node order, transition hashes, fixed-point/cycle counts and archived initial-restriction contrasts. The computation is finite and deterministic. Binary Boolean dynamics, the chosen external signals, uniform starting measure and wg-only fixed-point endpoint remain modeling choices. There is no physical intervention or wet-lab verification.

9.3 An exact reference-state check separates marker retention from state recovery

The previous maintained-clamp census deliberately used $wg=1$ at a fixed point as its endpoint. A new frozen-plan audit asks which complete fixed point is reached instead. Under the unrestricted condition $hh\ external=1$, $WG\ external=0$, $SLP=0$, the original model has one wg -ON fixed point and exactly 128 of 16,384 starts reach it. We identify it by every one of the fourteen free coordinates, not merely by a fate label. Its ON coordinates are CIA, CI protein, PH, PTC protein, SMO, WG protein, ci, ptc and wg; CIR, EN protein, HH protein, en and hh are OFF. The named software reference is not an experimentally established wild-type phenotype.

For each of the 28 maintained interventions, all 8,192 starts already matching the forced value are classified by their exact terminal fixed-point integer or by a cycle label. We score both complete reference equality and equality only on the thirteen nonmanipulated coordinates. The latter deliberately permits a disagreement at the forced node, while exposing downstream mismatches that marker wg cannot detect. A clamped value opposite to the original reference makes full-state equality impossible by construction. It need not make projected equality impossible. These tests answer different questions and must not be mixed in an intervention ranking.

Table 14: Selected exact maintained-clamp endpoints, percentages of the same reduced starting space. Projection omits only the manipulated coordinate. All values are model calculations, not measured rescue rates.

Forced condition	wg marker	thirteen-coordinate reference	full reference
CIA ON	25.000	25.000	25.000
PTC protein OFF	0.781	0.000	0.000
hh ON	0.781	0.000	0.000
ptc OFF	0.781	0.000	0.000
HH protein ON	0.781	0.781	0.000
PH OFF	0.781	0.781	0.000
wg ON	100.000	100.000	100.000

Three interventions make the endpoint distinction especially clear. PTC protein OFF, hh ON and ptc OFF each preserve wg -ON fixed-point outcomes from 0.78125% of their starting spaces, yet none matches the original reference even on the other thirteen coordinates. For them the disagreement propagates beyond the forced node. Calling these outcomes wild-type retention based on wg alone would conceal a change in the model's remaining expression pattern. By contrast, HH protein ON, PH OFF, SMO OFF and WG protein OFF retain projected-reference outcomes at 0.78125%, but cannot satisfy whole-state equality because their forced coordinate itself contradicts

the original reference. The projection explicitly states what is being ignored rather than quietly relaxing a phenotype.

CIA ON reaches the exact complete reference from 25% of starts, CI protein ON from 12.5%, ci ON from 6.25%, and EN protein OFF from 3.125%. Those compatible interventions change basin counts without changing the model's final reference pattern. Forcing wg ON does more than satisfy its own marker here: all reduced-space starts reach the complete original reference under this transition. That separate identity check removes the marker-only tautology, but it still does not show physical rescue. The intervention forces the initiating coordinate forever, the starting measure is artificial and uniform, and the external condition is fixed. It is neither recovery after a temporary perturbation nor robustness of an unforced tissue.

Graph classification is checked against independent scalar maintained trajectories over all reduced-space starts for CIA ON, PTC protein OFF and wg ON. The old marker fractions and all historical files remain untouched. The source hash, exact original reference, all fixed endpoints, graph hashes and all 28 conditions are in `results/grn_full_reference_audit.json`, generated by `experiments/grn_full_reference_audit.py`. The corrected conclusion is narrow: even exact model intervention counts depend on which state coordinates define success. Identifying a biological phenotype, measuring intervention kinetics and testing transient rather than permanent forcing remain separate research steps.

9.4 Transient forcing followed by release changes the intervention question

Permanent forcing changes the transition graph indefinitely. It cannot establish that a short intervention returns a system to the original attractor once the forcing ends. A new plan therefore evaluates finite forcing and explicit release. We hold the original external condition and the exact fourteen-coordinate reference fixed, start from every one of the 16,384 unrestricted states with equal weight, overwrite the chosen coordinate at intervention onset, maintain it for a stated number of synchronous updates, then restore the original successor without any clamp. Each released state is classified by its eventual original attractor. Thus success is attraction after release, not wg occupancy while it is forced.

The grid uses all fourteen nodes, both binary values and durations one, two, four and eight, giving 112 finite interventions. Every condition shares the full 16,384-start denominator, so comparisons do not confuse a reduced-state conditional measure with a different population size. Overwriting the coordinate initially maps pairs of starts onto the same clamped state, but both original starts keep their original weight. This initial overwrite is part of the hypothetical intervention. The durations count synchronous updates, not measured minutes, protein lifetimes or developmental stages.

Table 15: Exact original-reference attraction after release, percentage of the same 16,384 starts. No forcing remains during the classified return trajectory.

Transient forcing	one update	two	four	eight
CIA ON	3.125	6.250	25.000	25.000
CI protein ON	3.125	6.250	12.500	12.500
ci ON	3.125	6.250	6.250	6.250
EN protein OFF	3.125	3.125	3.125	3.125
wg ON	6.250	12.500	50.000	100.000
PTC protein OFF	.781	.781	.781	.781

With no intervention, 128 starts, or .78125%, reach the reference. CIA ON for one update gives 3.125% after release, two updates 6.25%, and four or eight updates 25%. CI protein ON saturates at 12.5% in the recorded four/eight-update conditions, while ci ON saturates at 6.25%. EN protein OFF gives 3.125% at every tested duration. These finite differences show dependence on intervention duration and node identity within this model; they do not identify a medically or developmentally useful dose. The complete grid retains OFF conditions with zero recovery rather than presenting only favorable interventions.

wg ON rises from 6.25% after one forced update to 12.5%, 50% and 100% after two, four and eight updates. The eight-update result is checked after the clamp has ended: all starts enter the original reference basin. This is stronger than the marker tautology of permanently forcing wg, but it remains a result of the stated single-cell binary rules and fixed external inputs. A finite-pulse biological claim would require intervention kinetics, coordinate meanings, measured starting-state distributions and phenotype validation. The code supplies none of those measurements.

PTC protein OFF illustrates another distinction. Under maintained forcing it could retain wg marker states that differed from the reference on other coordinates. Under finite forcing and release, the original dynamics recovers the reference from .78125% of full starts at every recorded duration. This is a restoration of the original rule, not evidence that permanent knockouts are equivalent to temporary perturbations. Exact reference equality at the instant of release is recorded separately from eventual reference attraction: a state outside the reference can still lie in its basin. Counting only occupancy at release would conflate immediate expression matching with later model recovery.

The experiment retains every release-state count, the number of distinct released states, immediate reference occupancy and eventual reference counts in `results/grn_release_audit.json`. Its frozen plan and executable are `experiments/grn_release_plan.json` and `experiments/grn_release_audit.py`. Tests regenerate the exact graph-based grid and independently enumerate full scalar trajectories for wg ON for two updates and CIA ON for four updates. All prior maintained/initial-restriction/noisy result files remain intact. This adds a controlled temporal-intervention

question, not a universal robustness law. Temporary forcing, permanent forcing, initial-value restriction and continuing random noise are four different processes and must not share an unqualified “basin rescue” number.

9.5 Initial-measure sensitivity is separate from endpoint identity

The uniform 16,384-state census is an exact counting measure, not a measured distribution of embryonic starting states. A further frozen audit keeps the original model, synchronous update and eventual exact-reference endpoint fixed while changing only the initial-state weights. We test independent bits with ON probabilities .25 and .75, and a uniform distribution on the 106 states within Hamming distance two of the original reference. The latter is deliberately reference-centered and cannot be treated as an unbiased embryo sample. Ten predeclared forcing conditions plus the unforced baseline are evaluated under all four measures. No measure is selected after its result.

Table 16: Exact-reference probability after release under artificial initial measures. Fractions are weighted census sums, not biological probabilities or sample estimates. The near-reference support has 106 states; all other measures have 16,384.

Condition	uniform	bit .25	bit .75	near reference
Unforced	.00781	.00165	.00494	.27358
wg ON, 2 updates	.12500	.14063	.04688	.63208
wg ON, 8 updates	1.00000	1.00000	1.00000	1.00000
CIA ON, 4 updates	.25000	.18750	.18750	.74528
PTC protein OFF, 1 update	.00781	.00165	.00494	.27358
EN protein OFF, 1 update	.03125	.00293	.07910	.43396

The unforced probability changes from .78125 percent under uniform bits to .16479 and .49438 percent under the two biased-bit measures, and to 27.35849 percent under the reference-centered measure. The model and trajectories have not changed: only the probability assigned to the same starts changes. In contrast, eight updates of artificial wg forcing reach the original reference from every enumerated start, so this particular model result is invariant to these weights. That invariance is an all-state property of this rule and forcing protocol, not evidence of physical rescue or general developmental robustness.

EN-protein OFF illustrates why fold comparisons depend on the measure. Its one-update probability is 3.125 percent under uniform weighting but .29297 percent under bit-.25 and 7.91016 percent under bit-.75. The unforced denominators also change. A quoted fold gain without the distribution would conflate an intervention response with the chosen prior. PTC-protein OFF remains identical to the unforced baseline under every tested measure. These findings preserve the

full-reference and projected-reference distinctions from permanent forcing.

For every tested released condition, the original dynamics yield no cycles and every wg-ON fixed endpoint is the same exact reference. Thus the two endpoint definitions coincide here after release; this is a measured property of these conditions, not a reason to replace exact identity with a marker in maintained-forcing experiments. The complete 44 weighted rows, support counts, source hash and plan hash are stored in `results/grn_initial_measure_audit.json`. Tests reproduce the full census, match uniform rows to the release audit and verify the 29-of-106 near-reference baseline. These artificial weights expose dependence on an assumption; empirical embryo distributions, asynchronous timing, altered rules and independent biological endpoints remain unvalidated.

9.6 External-input regimes invalidate an unconditional reference-rescue claim

The main forcing experiments hold hh-external ON, WG-external OFF and SLP OFF. A frozen context audit now enumerates every binary combination of those three external inputs while preserving all fourteen internal rules and synchronous timing. Each context receives the same 16,384 internal starts, the unforced baseline, eight forced updates of wg ON, and four forced updates of CIA ON followed by release. We retain the original internal reference as one endpoint, any wg-ON fixed state as another, and cycles separately. Before interpreting reference attraction, we test whether that original internal state remains a fixed point under the new context. Only three of eight

Table 17: Exact original-internal-reference attraction by external-input context. Every column has 16,384 uniform internal starts. Context order is hh-external, WG-external, SLP. Endpoint validity is reported rather than assuming a reference from one context remains meaningful everywhere.

Context	original reference fixed?	unforced	wg ON 8	CIA ON 4
000	no	0	0	0
001	no	0	0	0
010	no	0	0	0
011	no	0	0	0
100	yes	128	16384	4096
101	yes	16384	16384	16384
110	no	0	0	0
111	yes	16384	16384	16384

contexts preserve the original internal state as a fixed endpoint. The other five have zero attraction to it after release, not merely a smaller basin. In contexts 101 and 111 every unforced start already reaches the same internal reference, so neither forcing protocol supplies a gain there. In the original 100 context the previously reported counts reproduce exactly. Thus eight wg-forcing

updates yielding all-state attraction is a conditional result under a specified context, not robustness across external input settings. An internal-state match also does not claim equality of the external biological environment.

The wg-ON fixed endpoint count equals original-internal-reference attraction in all 24 tested rows and no cycle occurs. This agreement is an output of these exact censuses; it cannot validate the earlier maintained-clamp marker by analogy. Other contexts can make the prior endpoint invalid even if a different internally or biologically defined phenotype could be suitable. We do not select a new preferred reference after viewing outcomes. The complete fixed-endpoint lists, validity flags, source hash, plan hash and counts are in `results/grn_external_regime_audit.json`. Tests enumerate all eight contexts and reproduce the original release counts. This is context sensitivity of a published binary rule, not altered-rule, asynchronous, kinetic, empirical-prior or biological model validation.

9.7 Deterministic sequential update orders alter the basin

Synchronous timing is another assumption, separate from external context and initial-state weights. A frozen audit uses the original 100 external context and all 16,384 internal starts but updates the fourteen nodes in place through a complete deterministic sweep. Three orders are retained: source order, reversed source order and a fixed permutation from seed 2183. A forced node is overwritten initially and after every microstep, then released after the stated number of complete sweeps. The original reference is still a fixed point in every order. These are not random asynchronous trajectories, nor do sweeps have a physical duration equivalent to synchronous steps. The uniform

Table 18: Original-reference attraction under fixed sequential sweep schedules. Every cell counts 16,384 starts. The synchronous comparison uses updates, whereas sequential rows use complete fourteen-node sweeps; they do not represent paired physical pulse times.

Schedule	unforced	wg ON 8	CIA ON 4
Synchronous update, earlier audit	128	16384	4096
Source-order sweep	1024	16384	8192
Reverse-source-order sweep	1024	16384	4096
Fixed-permutation sweep	1024	16384	8192

unforced basin is 6.25 percent under each tested sequential schedule rather than .78125 percent under synchronous updates. That eightfold difference occurs without changing the rules or initial distribution. Equality across three sequential orders does not prove invariance across every order or random schedule. CIA forcing reaches half the starts for source/permutation orders but one quarter

for reverse order; wg forcing reaches all starts after eight sweeps in these three contexts. The audit does not identify a best biological update schedule or use a favorable schedule as the new default.

Complete-sweep terminal labels distinguish fixed endpoints from cycles of the sweep map. No cycle is observed in these nine conditions. Tests independently require every unforced sweep-fixed state to satisfy all original synchronous rules, so endpoint stasis here is not a hidden microstep excursion. Graph hashes, exact orders, counts and plan identity are retained in `results/grn_ordered_update_audit.json`. This experiment shows model-timing sensitivity and a limited all-state result for selected schedules. Stochastic asynchronous updates, schedules changing between sweeps, kinetic rate models and independently measured developmental trajectories remain open. No timing convention in this audit validates a physical pulse or embryo rescue.

9.8 Rule perturbations must pass endpoint-validity checks first

A frozen preliminary audit changes one syntactically named regulator in one target rule to constant zero or one, replacing all occurrences of that regulator within that rule. The 29 distinct regulator–target references yield 58 counterfactual models, including external-input references. The original 100 external context and original internal reference are fixed. This is a Boolean-expression counterfactual, not an experimentally measured edge deletion or gene knockout; constants zero and one represent different changed logical assumptions. In twenty models the original reference fails at the first

Table 19: Original-reference fixed-point validity before any counterfactual basin claim. The audit does not enumerate altered-model basins or select a new reference after observing invalidity.

Counterfactual models	58
Original internal reference remains fixed	38
Original internal reference is not fixed	20

successor check. A basin described as attraction to this fixed endpoint is therefore invalid for those models before any enumeration or rescue interpretation. In the other thirty-eight, passing one fixed-point check says nothing about basin size, stability or empirical correctness. Every changed coordinate is the perturbed target only, as expected for a one-rule successor audit. No altered-model attraction or biological result is claimed. The complete target, regulator, replacement value, changed coordinate and modified-model hash are retained in `results/grn_rule_validity_audit.json`; tests reproduce all 58 unique variants and the twenty invalid cases. This audit makes the reference-validity gate explicit rather than inheriting the original endpoint across changed models. Empirical benchmark scope still requires measured developmental states, assay definitions and an independent model; logical counterfactuals do not supply those observations.

10 Restricted-input comparisons for positional decoding

Verdict reframe 2 asks: not “how accurately can we decode” but “what is the minimum biological information required to recover developmental states”. We degrade the Bicoid gradient input to the positional decoder along three axes under the same disjoint-fly-line 5-fold protocol as the head-to-head benchmark, and ask where held-out error crosses the published Liu 2013 true-dose threshold model (4.94% EL).

Table 20: Information-degradation sweep (held-out RMSE, % egg length). Restricted-input scores are compared with the 4.94% EL true-dose threshold; this is not the best comparator in the archived benchmark.

Degradation	level	RMSE (%EL)
full gradient (64 points)	-	2.39
anterior window only	$x \leq 0.1$ EL	2.99
sparse sampling	2 points	2.53
sparse sampling	1 point	4.29
additive noise	$\sigma = 3 \times$ signal std	4.48
additive noise	$\sigma = 4 \times$ signal std	5.03 (crosses)
benchmark: Liu 2013 true-dose model		4.94

The single-point decoder reaches 4.29% EL and the anterior-window decoder 2.99% EL under the tested line-disjoint protocol. Both beat the 4.94% EL true-dose threshold comparator, but neither beats the measured-amplitude threshold comparator at 2.8566% EL in `results/cf_decoder.json`. The noise crossing is a property of the selected synthetic corruption and estimator, not a minimum biological information threshold. We did not measure information in bits, the embryo’s readout circuit, sensing noise or achievable optimal accuracy. Restricted-input prediction therefore does not establish that living morphogen readout needs no gradient shape. The full empirical sweep remains in `results/info_threshold.json`.

10.1 Update-rule and stopping-rule sensitivity: terminal wg status, not stochastic basins

The archived random-order update experiment uses all 16,384 starting states, eight seeded schedules and at most 300 sweeps. It stops early if no node changes during an entire sweep, then checks terminal `wg=1`. Its mean terminal fraction is 11.0413% with schedule SD 0.1226 percentage points, about 14.13 times the synchronous fixed-point-basin fraction 0.78125%. The output does not record how many trajectories hit the horizon without convergence, so it cannot certify eventual attraction

for every start. These numerical results establish sensitivity under a stated finite protocol, not a measured asynchronous developmental outcome.

The noisy experiment samples every fourth initial state, four repetitions, and scales counts by four. The subset is systematic rather than a random sample, and its count variation is not a population uncertainty interval. Each sweep first applies the synchronous rule, then independently flips each free node with probability q . It stops as soon as a noisy successor happens to equal the current state, or at 300 sweeps, and checks terminal wg status. With continuing noise, an unchanged step is not proof of permanent absorption. Terminal $wg=1$ is also not a test of the complete wild-type fixed point.

At $q=.001$ the archived terminal fraction is 0.5554%, below the synchronous 0.78125%. At .01, .05 and .1 it is 4.9255%, 18.1702% and 20.4834%. Thus the earlier wording that noise simply inflates the basin is false even within this finite library. The four recorded fractions and lower-noise decrease remain unchanged in `results/grn_update_rules.json`. They must not be relabeled stochastic basin volumes or biological robustness.

The original asynchronous and noise scripts do not recompute the per-node perturbation index or maintained-clamp interventions under each new rule. Consequently they do not demonstrate that the same CI/EN axis remains dominant across update schemes. Actual cellular timing, perturbation distributions and fate readout are unmeasured. The useful result is protocol dependence that motivates fresh, explicitly defined absorption or horizon-specific experiments, with failed claims preserved rather than repaired by biological analogy.

10.2 Paired continuing-noise paths separate the stopping rule from the endpoint

A new preplanned software audit holds each path prefix fixed while comparing first-unchanged stopping with continued noisy updating. Its frozen local plan uses four replicates, each with 4,096 uniformly sampled initial states (sampling with replacement), five noise rates and three horizons. The same initial states and random uniforms are shared across rates within a replicate. This pairing reduces comparisons to a specified protocol difference; it does not make the paths or horizons statistically independent. The plan was frozen before execution and its SHA-256 is retained in the result record. It is a local plan, not an externally registered study.

For each path, we store two endpoints. The marker endpoint asks only whether $wg=1$ at the observation time. The stricter endpoint asks whether the entire fourteen-node state exactly matches an original unrestricted synchronous wg -ON fixed point; this model has one such state under the chosen external condition. This is an explicit full-state software reference, not a biologically validated wild-type phenotype. With continued noise, that reference can be visited and then left. Even

a full-state match is therefore occupancy at a stated horizon, not permanent rescue.

The continuous process is $X_{t+1} = f(X_t) \oplus Z_t$, where each free-node bit of Z_t is independently Bernoulli(q). At $0 < q < 1$ every bit-mask has positive probability, so from any current state every possible next state has positive probability. No single state is permanently absorbing. An unchanged step has probability greater than zero, but says nothing about the next sweep. First-unchanged stopping selectively freezes such events and changes the process being measured. We do not infer the stationary distribution or claim that 1,000 steps are sufficient for equilibrium.

Table 21: Paired stopping/endpoint audit at 1,000 steps. Entries are mean percentages across four replicates of 4,096 random starts. SDs across replicates and all 100/300/1,000-step records are saved; these are not confidence intervals.

q	stopped wg	continued wg	stopped full state	continued full state
0	0.739	0.739	0.739	0.739
0.001	1.215	12.622	1.215	12.115
0.01	5.334	13.208	5.310	9.253
0.05	17.560	15.479	16.193	2.466
0.1	21.069	19.061	15.631	0.488

At $q = .001$, first-unchanged stopping gives a marker fraction of 1.215%, unchanged over the three recorded horizons because all sampled paths have stopped before 100 sweeps. Continuing the same prefixes instead gives 7.043%, 11.700% and 12.622% at 100, 300 and 1,000 sweeps. The large horizon effect was hidden by the old stopping rule. This does not contradict the preserved historical 0.5554% record: that record used a systematic quarter of the starting states and different random draws. The fresh random-start measure is deliberately separate, not a silent replacement.

The distortion is not uniformly upward or downward. At $q = .1$ and 1,000 steps the stopped marker fraction is 21.069% versus a continued fraction of 19.061%. More importantly, the full-reference fractions are 15.631% versus 0.488%. A marker-only reading would hide most of the loss of full-state occupancy under stronger continuing noise. Across the tested rates, continued wg occupancy rises while continued reference-state occupancy peaks at the low nonzero rate and then falls. These rates do not establish an optimized biological noise level: initial-state weighting, external condition, horizon and endpoint were fixed model choices, not measured developmental inputs.

The noiseless lookup audit reproduces the exact 128/16,384 basin count. Every vectorized transition-table entry is checked against the original scalar Boolean successor, and independent noiseless trajectories reproduce the saved random-start results. The simulation makes no intervention, changes no historical result, and supplies no wet-lab evidence. All 60 replicate/rate/horizon records, the original reference state, node order, model hash and frozen-plan hash are saved in

results/grn_noise_horizon_audit.json; execution is experiments/grn_noise_horizon_audit.py. This distinguishes a reproducible protocol failure from an unsupported developmental story.

10.3 A reference-valid rule change can still alter its exact basin

A separately frozen follow-up enumerates every internal start for all 58 one-rule constant replacements already audited for reference validity. The external context, synchronous convention and uniform measure over 16384 internal starts remain fixed. The endpoint is the exact original internal reference, not merely a wg-ON marker and not a newly selected preferred state. For the 20 variants in which that reference is not fixed, its basin endpoint is inadmissible and is recorded as null rather than as a comparable zero. Fixed and cyclic termination counts are still retained for every variant. Of the 38 reference-valid variants, 30 retain the baseline 128 exact-reference starts and eight enlarge that basin. Thus passing the reference-validity check does not establish unchanged dynamics or robustness. The complete eight changed valid variants are retained below; this selection is a result summary, not an exclusion from the exhaustive ledger.

Table 22: All reference-valid variants whose exact-reference basin differs from the baseline 128 of 16384 starts. Replacement is a mathematical constant in one target rule, not an empirical edge knockout.

Target	Regulator reference	Constant	Exact-reference starts
CIA	CI protein	1	2048
CI protein	ci	1	1024
EN protein	en	0	256
ci	EN protein	0	512
wg	CIA	1	4096
wg	CIR	0	256
wg	SLP	1	16384
wg	wg	1	16384

The two globally absorbing original-reference cases arise from constant forcing within the wg update rule, not a demonstrated biological rescue. Across all 58 variants, cyclic-start counts range from zero to 16384; two variants send every start to cyclic termination. A Boolean counterfactual can destroy the reference or change its basin without fitting any developmental measurement. These cases cannot be ranked as interventions in a living embryo. The original-model result, reference-invalid negatives and context dependence remain unchanged. The vectorized successor enumerator is independently checked against the existing scalar Boolean parser for every

state in every altered model, including negated constants. Tests also reproduce complete end-point aggregation, uniform denominators, 20 null inadmissible endpoints and the eight changed valid basins. The frozen plan and full ledger are `experiments/grn_rule_basin_plan.json` and `results/grn_rule_basin_audit.json`. This exact model calculation fills a basin-analysis gap while leaving empirical regulatory validation, spatial tissue structure, continuous dynamics and biological intervention inference open.

10.4 Local shell returns depend on the same rule counterfactuals

A further exploratory plan partitions all 16384 internal starts by Hamming distance zero through fourteen from the exact original reference. The rule variants and external context remain fixed. This conditional measure asks about nearby state perturbations rather than uniform whole-state volume, but it is mathematical weighting, not a measured developmental perturbation distribution. All 58 counterfactuals and all shells are retained; the 20 reference-invalid variants remain null rather than being assigned return probabilities. The baseline exact-reference basin has 128 starts but returns from seven of fourteen one-bit perturbations and twenty-one of ninety-one two-bit perturbations. Those local fractions, .5 and .23077, must not be replaced with the global .0078125 fraction. Of the 38 reference-valid counterfactuals, thirty preserve both shell counts and eight change them. The complete changed set is below.

Table 23: All reference-valid rule replacements changing one-bit shell return. Denominators are fourteen and ninety-one starts, respectively. Constant forcing is not a biological intervention result.

Target	Regulator reference	Constant	Radius one	Radius two
CIA	CI protein	1	11/14	55/91
CI protein	ci	1	10/14	45/91
EN protein	en	0	8/14	28/91
ci	EN protein	0	9/14	36/91
wg	CIA	1	12/14	66/91
wg	CIR	0	8/14	28/91
wg	SLP	1	14/14	91/91
wg	wg	1	14/14	91/91

The valid-variant one-bit return range is .5 to one and the two-bit range is .23077 to one. These larger counts agree with the previously recorded global basin enlargements; they are not an independent replication. Choosing a near-reference starting measure changes the numerical estimand, not the biological validity of the Boolean rules. The globally absorbing constant-wg cases remain mathematical forcing rather than tissue rescue. Every shell denominator equals its binomial coefficient,

and summing shell returns exactly reproduces each prior global basin count. Plans and complete shell counts remain in `experiments/grn_rule_radius_plan.json` and `results/grn_rule_radius_audit.json`. Empirical initial-state distributions, regulatory perturbations and independent physiological validation remain unmeasured.

10.5 Finite-step return latency is distinct from eventual basin membership

A frozen follow-up at head 9b27ea8 computes the exact first hitting time of the original reference for every one of the 16,384 states in the baseline and all 58 previously specified single-rule constant counterfactuals. The original external context and synchronous update convention are unchanged. The 20 variants for which the reference is not a fixed point retain null latency summaries, rather than an invented return endpoint. For each of the 38 admissible variants, starts are partitioned into the same exact Hamming shells as the preceding basin audit. Histograms retain each finite first-hit time; cumulative horizons are fixed in advance at 0, 1, 2, 4, 8, 16 and 32 steps.

The baseline’s eventual return counts at radii one and two are seven of fourteen and 21 of 91. These are not immediate returns: after one step the counts are four and six, and after two steps six and fifteen. Every returning baseline start arrives by step three. Across valid rule counterfactuals, maximum first-hit times range from two to eight steps, with observed maxima in the set $\{2, 3, 4, 5, 8\}$. Thus an eventual basin comparison conceals timing differences under exactly the same finite model.

Radius	Horizon	Baseline count	Valid-variant range	Changed variants
1	1	4	4–5	7
1	2	6	6–8	8
1	4	7	7–10	8
2	1	6	6–10	7
2	2	15	15–28	8
2	4	21	21–45	8

Each count uses the complete shell denominator, fourteen at radius one and 91 at radius two, not only eventual returning starts. The table’s changed-variant count compares each finite-horizon count with the baseline at that horizon. The full histograms sum exactly to all previously saved shell-basin counts. Tests independently traverse every functional graph on up to four states and compare every fixed endpoint’s first-hit times, including starts in other fixed basins or nontrivial cycles. The numerical follow-up was selected after shell-basin outcomes, not as an independent replication.

Discrete synchronous model steps are not measured developmental time. Constant substitutions of a regulator within one rule are not biological knockouts. Neither shorter latency nor a larger

mathematical basin establishes rescue, robustness in tissue, or independent physiological validation. All inadmissible endpoints, delayed returns and nonreturning starts are preserved in `results/grn_rule_latency_audit.json`, with the frozen plan in `experiments/grn_rule_latency_plan.json` and the runner in `experiments/grn_rule_latency_audit.py`. The manuscript content-page gate remains uncertified.

10.6 Basin boundary-to-volume fragility under the same rule variants

A frozen follow-up at head 6821374 asks whether the previously observed enlarged basins also change a distinct exact statistic: the fraction of all directed single-internal-bit flips from basin members that leave the original-reference basin. The original external context, synchronous successor rules and all 58 constant-substitution variants are unchanged. Every valid model is enumerated over all 16,384 states, and every one of fourteen node flips is counted for every basin member. The 20 inadmissible reference endpoints retain null boundary summaries rather than zero fragility. The pre-registration SHA256 is `f13be7fe4e43aae2802547aa8d8ac63063af857192d357eb62b495a4ca6efab8`.

The baseline basin has 128 members, 1,792 directed flips and 896 exits, giving exit fraction .5. Seven nodes always exit this basin under a single flip and seven never do. These seven exit-sensitive nodes also match the reference-only flip exits in this baseline. The statistic weights all basin members and all nodes uniformly; it is not a measured perturbation probability.

Target	Replaced regulator/value	Basin	Directed flips	Exits	Fraction
CIA	CI protein/1	2048	28672	6144	.21429
CI protein	ci/1	1024	14336	4096	.28571
EN protein	en/0	256	3584	1536	.42857
ci	EN protein/0	512	7168	2560	.35714
wg	CIA/1	4096	57344	8192	.14286
wg	CIR/0	256	3584	1536	.42857
wg	SLP/1	16384	229376	0	0
wg	wg/1	16384	229376	0	0

These are all eight admissible variants whose basins differ from the baseline. Their exit fractions are lower; the remaining thirty valid variants retain fraction .5. Absolute exit counts need not fall: the CIA-input substitution gives 6,144 exits rather than 896 while the denominator grows enough to lower the fraction. Hence basin size, exit count and exit fraction are separate quantities and should not be called interchangeable robustness evidence. Full-state absorbing cases have no boundary by construction, not demonstrated biological rescue.

All per-node counts and reference-only exits remain in `results/grn_rule_boundary_audit.json`. Tests reconcile every basin size with the existing exact basin ledger, check directed flip accounting and independently traverse every four-state functional graph to verify its fixed-reference boundary. The plan and runner are `experiments/grn_rule_boundary_plan.json` and `experiments/grn_rule_boundary_audit.py`. This follow-up was selected after the shell and latency results and is not independent replication. Uniform basin-state/node weights, fixed external conditions and constant rule replacements are mathematical choices, not empirical developmental noise or biological knockouts. Spatial, continuous and physiological validation remain open; the content-page gate is uncertified.

10.7 Continuous equations change the Boolean rule, and the feedback threshold depends on parameters

The reduced five-node Boolean core and the Hill ODE share a signed-edge list, not a transition semantics. The Boolean successor turns a node on for positive signed input, off for negative input, and preserves its current value when the signed sum is zero. This tie rule gives persistence without an explicit self-edge. In particular, $(wg, en, hh, ptc, ci) = (1, 1, 1, 0, 0)$ is a Boolean fixed point despite $ci=0$. The ODE instead has linear decay and no production without activation: at that same state, without added feedback, $wg = -1$. Removing zero-input persistence while introducing a continuous production function changes two assumptions at once. A differing outcome does not isolate discreteness or show that real segment-polarity bistability requires a particular biological feedback strength. This five-node reduction is also not the fourteen-node public network used for the exact basin census above.

The old records in `results/ode_vs_boolean.json` and `results/ode_vs_boolean_scan.json` are retained as historical finite-grid terminal calculations. A finite set of initial values ending near one rounded state does not prove a single global attractor. The current tracked base experiment specifies an Euler integration, clipping and approximate rounded final-state grouping; it does not certify stationary points by their residuals or prove their stability. The threshold-scan JSON is not accompanied by a tracked scan-generating script specifying its added autoactivation term and parameter cross-product. Its unqualified reading that every tested parameter pair requires $a \geq .9$ is therefore withdrawn, as is the biological near-full-strength requirement. No archived number is changed.

To make the comparison inspectable, a fresh local plan defines a separate continuous model

explicitly. Let $H(x) = x^n / (K^n + x^n)$. Its five derivatives are

$$\begin{aligned} \dot{wg} &= \min(1, H(ci) + aH(wg)) - wg, \\ \dot{en} &= H(wg) - en, \quad \dot{hh} = H(en) - hh, \\ \dot{ptc} &= H(ci)(1 - H(hh)) - ptc, \quad \dot{ci} = -ci. \end{aligned}$$

The last equation follows the production convention that a node with no positive activator has zero production. It shows directly that this continuous model cannot maintain positive ci : that is a modeling choice, not an observed biological loss of expression. An added wg term creates a different feedback system; it is not a calibrated strength of cis -regulation. Each term, saturation convention and decay rate must be stated before a threshold means anything.

The frozen plan scans four K values, three Hill exponents, five added strengths ($a = 0, .5, .85, .9, 1$), and all $3^5 = 243$ starts on $\{0, .5, 1\}^5$. We integrate to $t = 100$ using unclipped fourth-order Runge–Kutta at step sizes $.1$ and $.05$. The 120 parameter/step-size records count terminal wg above $.5$ and record full-vector residuals separately. The largest endpoint difference after halving the step is 1.875×10^{-8} , small for this protocol. The largest absolute derivative at the horizon is 1.015×10^{-4} , so we do not label every terminal vector a certified equilibrium. Rounded endpoint counts remain descriptive rather than attractor counts.

Table 24: New explicitly defined Hill model: percentage of 243 grid starts with wg above $.5$ at $t = 100$, RK4 step $.05$. The parameter dependence contradicts a universal near-full-strength threshold; entries are finite-grid outcomes, not continuous basin volumes.

K	n	$a = .85$	$a = 1$
0.2	2	88.889	88.889
0.2	4	88.889	88.889
0.2	8	88.889	88.889
0.35	2	88.889	88.889
0.35	4	88.889	88.889
0.35	8	77.778	77.778
0.5	2	0.000	66.667
0.5	4	0.000	66.667
0.5	8	55.556	66.667
0.65	2	0.000	0.000
0.65	4	0.000	0.000
0.65	8	0.000	44.444

Without added feedback ($a = 0$) the terminal wg -ON fraction is zero at every sampled setting, consistent with the earlier qualitative negative but not a global-attraction proof. At $K = .2$, $a = .85$

already gives 88.889% for each tested exponent. At $K = .65$ with $n = 2$ or 4, even $a = 1$ gives zero. At $K = .5, n = 2$, the $a = 1$ fraction is 66.667%, but its largest derivative residual is the audit maximum, so a near-threshold endpoint should not be promoted to a proven stable attractor. The new term is explicitly stated and not asserted to reconstruct the undocumented old threshold scan. These results demonstrate model and parameter dependence; they neither fit actual cellular rates nor validate a developmental intervention.

All settings, endpoint discrepancies, residuals, coordinate ranges and rounded endpoint counts are in `results/grn_ode_semantics_audit.json`, generated by `experiments/grn_ode_semantics_audit.py`, with frozen-plan hash. Tests verify the Boolean persistence counterexample, derivatives at the same state, solver step-halving and representative saved outcomes. A more meaningful biological comparison would require a justified mapping between rules and kinetics, a measured phenotype and parameters, and stability/continuation analysis beyond this finite grid. That evidence is absent and remains an open gap.

10.8 Four-model comparison: the metric and ridge intercept change the reading

The historical comparison on 171 cached Turing simulations reports linear-theory relative RMSE .3427, dispersion-curve CNN .3386, parameter MLP .4487 and curve-feature ridge 1.2396. These exact values remain in `results/turing_4model.json`. The .0041 CNN difference is a small point-estimate difference, not an established statistically meaningless effect or a demonstrated improvement: the saved summary does not include paired predictions, intervals or a hypothesis test. Neither a small edge nor an unsuccessful learner proves what information its inputs contain. We withdraw the earlier claims that only linear-theory features matter and that the shift is unlearnable from those views.

The target is $y = (k_{sim}^2 - k_{min}^2)/(k_{max}^2 - k_{min}^2)$. The recorded relative RMSE is

$$\sqrt{N^{-1} \sum_i ((\hat{y}_i - y_i)/(|y_i| + 10^{-9}))^2},$$

which weights squared target errors by inverse squared target magnitude. This differs from ordinary target RMSE. A target near zero receives much more weight; changing the metric can change the ranking without changing a single prediction. Further, the historical fixed linear predictor is scored globally on all starts, while the trained learners are summarized by the arithmetic mean of twelve fold RMSEs (six folds, two partition seeds). Mean fold root errors and a pooled root error need not agree. Two partitions of the same 171 simulations are not two independent physical replications.

The ridge implementation also standardizes its features within each training fold but fits no intercept. Centered features cannot represent the nonzero mean target using a linear slope alone.

That is a comparator specification defect, not evidence that curve summaries add noise rather than signal. A fresh frozen local audit retains identical folds, features, ridge penalty one and targets, and adds an unpenalized intercept using only the training-fold target mean. It includes the original no-intercept version, the fixed linear predictor and a training-mean constant. No CNN or MLP is retrained, and no hyperparameter is selected using these outcomes.

Table 25: Same saved simulation targets and folds: mean over the two partition seeds. Each model has out-of-fold predictions for all 171 simulations per seed. These are descriptive metric comparisons, not independent-replicate uncertainty.

Predictor	pooled rel. RMSE	fold-mean rel. RMSE	pooled RMSE
Linear theory	.3427	.3417	.1865
Training-mean constant	.7381	.7021	.1756
Ridge, no intercept	1.2466	1.2396	.4390
Ridge, with intercept	.3579	.3511	.1372

Adding the intercept reduces mean-fold relative ridge error from 1.2396 to .3511. Under that metric, the repaired ridge still does not beat the fold-matched linear value .3417. Under ordinary pooled target RMSE the same repaired ridge is better (.1372 versus .1865), and even the training-mean constant is better than linear (.1756). Thus a universal “nothing beats linear theory” heading is false once the loss is stated differently. This is a useful comparator audit, not a new optimized predictor or an externally validated performance gain. The training loss in the historical neural experiment is ordinary squared error, while its reported score is relative error; this mismatch is another reason not to turn its ranking into an information boundary.

The saved source SHA-256, local plan hash, all per-seed metric records, target range and mean are in `results/turing_metric_audit.json`, generated by `experiments/turing_metric_audit.py`. Tests reproduce the saved no-intercept result, verify the metric-dependent ranking and regenerate this deterministic audit byte-for-byte. The historical result file and its negative interpretation remain visible as superseded claims, not silently edited measurements. Independent simulations, preselected loss, matched model selection and saved paired neural predictions would be needed for a broader comparison. Nothing here establishes how biological tissues predict or select wavenumber.

10.9 Cross-model contrast does not isolate a topology law

Verdict #6 asks whether the fragility framework applies to a second developmental network. We applied the same basin-enumeration and single-bit-exit construction to the four-node gap-gene cascade with an anterior Bcd prepatter. This implemented module has one global attractor: its WT

basin fraction is 1.0 and its basin-average one-bit exit fraction is zero. The implemented segment-polarity network has a WT basin fraction of 0.781% and exit fraction 0.500. The numerical contrast is reproducible, but two different models do not isolate topology from their regulatory rules, external clamps, state definitions or update schedules.

We therefore withdraw the earlier claim that fragility is a property of feedback topology rather than modeling choices. The floral experiments in this paper supply direct counterexamples to that wording: with fixed topology, fate rankings change with starting measure, asynchronous update and recovery horizon. A single-global-attractor system has zero eventual basin-exit fragility by definition under its deterministic rule, not because every cascade or real developmental network is globally robust. Feedback circuits can have several attractors, but these examples do not prove that feedback necessarily causes a given degree of fragility.

The useful transferable object is the specified measurement protocol: identify the transition rule, state measure, perturbation kernel and horizon, then compare robustness under controlled changes. The two-model contrast motivates a topology hypothesis, not a validated causal law. Original numerical records remain in `results/grn_fragility_gapgene.json` and `experiments/fragility_gapgene.py`; their scope is unchanged.

11 Conclusion

Three digital-embryo studies, three different relationships between computation and data. Study A found a robust phenomenon (the beyond-onset wavenumber upshift, 15–60%, mid-band peak, resolution-verified) and an honest dead end (its CNN predictability, retracted at $n = 171$). Study B turned a curated Boolean model into exact numbers: a full attractor census under eight conditions, a wild-type basin of 0.77% of state space, and a basin boundary learnable from ~ 60 labelled states — plus the sampling degeneracy that makes naive versions of that measurement meaningless. Study C reproduced a canonical morphogen measurement on 676 real gradients, refined the dosage–length-scale coupling and localised it to the anterior embryo, broke the published cephalic-furrow decoding benchmark on its own data (2.38% vs. 4.94% EL, paired CI [1.99, 3.11]), and measured the dosage-dependence of positional decodability in gap genes (9.91 $\rightarrow$ 6.21% EL, CI [3.00, 4.40]).

The biological throughline comes first (verdict #8: every computational result must answer “what biological interpretation follows”). Study A: the tested CNN uses the linear dispersion curve, not pattern images, and fails to improve the held-out wavenumber prediction. The experiment does not establish an information-theoretic limit on local image features or a biological mechanism of non-local encoding. Study B: the wild-type segment-polarity state is a fragile attractor (fragility index $F = 0.500$) whose synchronous perturbation sensitivity is concentrated in seven nodes including *wg*, *ci* and *en*. The separate finite-horizon terminal-*wg* calculation changes with update

protocol - the historical clamp-labeled experiment only restricts initial node values, so its knockout interpretation is withdrawn. The terminal asynchronous/noisy wg fractions also depend on horizon and early stopping, not an eventual stochastic basin. Floral-state audits further show that robustness rankings depend on starting measure, update schedule and recovery horizon. These computational interventions do not establish how actual development maintains robustness. Study C: a supervised decoder can predict a measured landmark from restricted gradient inputs under this dataset and splitting protocol. The single-point score beats the true-dose 4.94% EL comparator but not the measured-amplitude 2.8566% EL comparator. The result does not identify whether a living embryo is limited by sensing, integration or another process.

The methodological throughline: every number lives in a committed results file produced by tested code from accession-level public data; implemented formulas and tests provide local checks, not a certification of every mathematical statement; failed experimental claims are retained alongside their corrections. We consider the retraction machinery — not any single result — the project’s most reusable output.

If there is one transferable claim, it is procedural: specific published-data benchmarks and model calculations can be reproduced and compared under stated splits with commodity hardware. Line-disjoint evaluation and paired resampling do not alone establish external validity, causal biological interpretation or universal superiority. The repository that backs this report — code, tests, manifests, results, and this text — is constructed so that any single number can be traced to the command that produced it in under a minute.

References

- [1] A. M. Turing, “The chemical basis of morphogenesis,” *Phil. Trans. R. Soc. Lond. B* **237**, 37–72 (1952).
- [2] A. Gierer and H. Meinhardt, “A theory of biological pattern formation,” *Kybernetik* **12**, 30–39 (1972).
- [3] L. Wolpert, “Positional information and the spatial pattern of cellular differentiation,” *J. Theor. Biol.* **25**, 1–47 (1969).
- [4] T. Gregor, E. F. Wieschaus, A. P. McGregor, W. Bialek and D. W. Tank, “Stability and nuclear dynamics of the Bicoid morphogen gradient,” *Cell* **130**, 153–164 (2007).
- [5] F. Liu, A. H. Morrison and T. Gregor, “Dynamic interpretation of maternal inputs by the *Drosophila* segmentation gene network,” *Proc. Natl. Acad. Sci. USA* **110**, 6724–6729 (2013). PMID 23580621; data: Zenodo record 4942019 / Dryad doi:10.5061/dryad.2j9p4.

- [6] M. D. Petkova, G. Tkačik, W. Bialek, E. F. Wieschaus and T. Gregor, “Optimal decoding of cellular identities in a genetic network,” *Cell* **176**, 844–855 (2019).
- [7] J. O. Dubuis, G. Tkačik, E. F. Wieschaus, T. Gregor and W. Bialek, “Positional information, in bits,” *Proc. Natl. Acad. Sci. USA* **110**, 16301–16308 (2013).
- [8] G. von Dassow, E. Meir, E. M. Munro and G. M. Odell, “The segment polarity network is a robust developmental module,” *Nature* **406**, 188–192 (2000). (BioModels BIOMD0000001065.)
- [9] R. Albert and H. G. Othmer, “The topology of the regulatory interactions predicts the expression pattern of the segment polarity genes in *Drosophila melanogaster*,” *J. Theor. Biol.* **223**, 1–18 (2003).
- [10] T. Helikar et al., “The Cell Collective: toward an open and collaborative approach to systems biology,” *BMC Syst. Biol.* **6**, 96 (2012). Models id-021, id-027 via the sybila/biodivine-boolean-models curated mirror.
- [11] B. Derrida and Y. Pomeau, “Random networks of automata: a simple annealed approximation,” *Europhys. Lett.* **1**, 45–49 (1986).
- [12] A. V. Spirov et al., SDD-model critique of Bicoid gradient formation, *PLoS ONE* (2017); data via PLOS figshare.

A Tool and dataset manifest

A.1 External tools and resources (40/40 gate)

Counting rule (strict standard): only research and data tools count toward the gate; development and build infrastructure (Git, LaTeX, pip, etc.) is listed in the manifest but deliberately excluded from the count. Every entry was genuinely used in producing a result, figure, lookup, or artifact of this project. Table 26 reproduces the committed manifest (`data/TOOLS AND DATASETS.md`).

A.2 Dataset manifest (accession-level)

Uniform rule: one identifier-backed record = one dataset.

The 29 fly lines of the LiveImaging modality, with their maternal *bcd* dosages as used in the DCLS analysis: 1IIA, 1IIC, 1XA, 2IIA, 2IIB, 2XA (1×); 1XA1IIA, 1IIA1IIIA, 2IIIA, 2IIIB, 2IIC, 1XA2IIA, 2XA1IIA (2×); 2XA2IIA, 2IIA2IIIA, 2XA2IIIA, 2XA2IIIB (3×); 2XA1IIA2IIIA, 2XA2IIA2IIIA,

Table 26: External tools/resources genuinely used (committed manifest).

#	Tool/resource	Used for
1	Zenodo (record 4942019)	Liu 2013 dataset download + API metadata verification
2	Dryad (doi:10.5061/dryad.2j9p4)	dataset provenance + Dryad API record
3	Cell Collective	published GRN id-021/id-027
4	sybila/biodyn-boolean-models	curated .bnet model files
5	BioModels (BIOMD0000001065)	von Dassow 2000 reference
6	NCBI eutils	gene-ID + citation lookups (data/lookups/ncbi_genes.json)
7	NCBI Gene	bcd 40830, hb 41032, eve 36039, kni 40287, gt 31227
8	PubMed	PMID 23580621 verification
9	PNAS/PMC (Liu 2013 + SI)	published bars: lambda, Sx, Sc
10	Gregor et al. 2007 Cell	gradient-precision benchmark
11	Petkova et al. 2019 Cell	4-gene decoding benchmark
12	Dubuis et al. 2013 PNAS	single-gene information benchmark
13	PLOS figshare (Spirov 2017)	SDD-critique cross-check
14	UniProt REST	protein records: bcd P09081, hb P05084, eve P06602, kni P10734, gt P39572
15	KEGG REST	dme04310 (Wnt), dme04341 (Hedgehog) pathway checks
16	CrossRef	Liu 2013 DOI 10.1073/pnas.1220912110 verification
17	OpenAlex	citation metadata (W2007062550)
18	Ensembl REST	bcd gene lookup (FBgn0000166)
19	InterPro (EBI)	bcd homeodomain annotation (IPR001356)
20	MyGene.info	cross-DB gene identifier reconciliation
21	STRING API	segment-polarity interaction network (141 edges)
22	RCSB PDB search API	Bicoid homeodomain structures (9RW7-9, 1ZQ3)
23	Europe PMC REST	literature metadata + citation count
24	EBI OLS4	GO term resolution (GO:0000978)
25	Gene Ontology API (AmiGO)	GO:0003700 term verification
26	Semantic Scholar Graph API	citation cross-check
27	NumPy	all numerics
28	SciPy	fitting, statistics, .mat I/O
29	pandas	tables
30	PyTorch (CPU)	CNN/GNN/MLP models
31	scikit-learn	decoder cross-checks
32	NetworkX	GRN graph handling
33	matplotlib	all figures
34	h5py	.mat v7.3 fallback
35	statsmodels	DCLS robust OLS (results/dcls_ols.json)
36	scikit-image	Turing morphometrics (re- sults/turing_morphometrics.json)

Table 27: Dataset counts. Total 2,858 (gate: 120).

Source	Records	Count
Liu 2013 LiveImaging (Bcd-GFP, 29 lines)	per-embryo gradients	676
Liu 2013 Immunofluorescence (Hb/Kr/Gt/Kni/Eve)	per-embryo records	2,180
Cell Collective id-021	model	1
Cell Collective id-027	model	1
Total		2,858

2XA2IIC2IIIA (4×); plus wild-type and un-annotated controls used only where dosage is not required.

B Exact attractor records (Study B)

The complete exact census from `results/grn_real.json`. Node order in the bit strings: CIA, CIR, CI, EN, HH, PH, PTC, SMO, WG-protein, ci, en, hh, ptc, wg (left to right; gene and protein nodes prefixed `v_` in the record).

Table 28: All fixed points of Cell Collective id-021 under all eight external conditions (synchronous update; no limit cycles exist in any condition).

Condition (hh-WG-SLP)	fixed point	state vector
hh0-WG0-SLP0	1/1	01100010010000
hh0-WG0-SLP1	1/1	01100010010000
hh0-WG1-SLP0	1/2	00011001001100
hh0-WG1-SLP0	2/2	00011010001100
hh0-WG1-SLP1	1/1	01100010010000
hh1-WG0-SLP0	1/2	10100111010010
hh1-WG0-SLP0	2/2	10100111110011
hh1-WG0-SLP1	1/1	10100111110011
hh1-WG1-SLP0	1/1	00011001001100
hh1-WG1-SLP1	1/1	10100111110011

Two readings. Under the null-signal conditions the unique attractor is the CIR/CI repressor state (01100010001000): *ci* is broadly expressed, its repressor fragment CIR dominates, and the segment-polarity genes *en*, *hh*, *wg* are all OFF. The wild-type bistable condition (hh1-WG0-SLP0) adds the 01001110010011 state in which *en*, *hh*, *wg* and their proteins are ON together with SMO

signalling — the *wg*-expressing stripe fate. SLP= 1 collapses the bistability in both directions, matching the published role of *sloppy paired* as a condition setter.

C Reproduction commands

```
git clone <repository-url> && cd mega27-14-digital-embryo
pip install -e . && pip install -r requirements.txt
python scripts/download_data.py          # re-fetch all public data
python -m pytest tests/ -q               # hermetic suite incl. sympy checks
python experiments/run_all.py            # regenerates results/results.json
# individual studies:
python experiments/turing_sweep.py       # Study A sweep (results/turing_sweep.json)
python experiments/turing_cnn.py         # Study A surrogate + retraction record
python experiments/turing_morphometrics.py
python experiments/grn_real.py           # Study B census
python experiments/grn_learning_curve.py
python experiments/grn_gnn_balanced.py
python experiments/bcd_sdd_test.py       # Study C fits (results/bcd_sdd.json)
python experiments/bcd_window_sensitivity.py
python experiments/dcls_ols.py
python experiments/cf_decoder.py         # benchmark (results/cf_decoder.json)
python experiments/gap_decode.py         # decoding (results/gap_decode.json)
python experiments/hb_boundaries.py
cd paper && lualatex main.tex            # this document (Times New Roman)
pdffonts main.pdf                       # font-embedding verification
```

Each script is deterministic given its seed set; the committed results files are the outputs of exactly these commands.

D Results-file inventory

Every quantitative claim in this report lives in one of these committed files (JSON; byte counts as committed):

Table 29: Committed results files and the claims they carry.

File	bytes	Carries
bcd_sdd.json	175,102	all 676 per-gradient fits; population λ ; per-dosage stats; F -test rejection fractions; flag-ship bootstrap CI
bcd_windows.json	1,133	window-sensitivity of λ and DCLS (Table 8)
dcls_ols.json	349	robust-OLS DCLS cross-check
cf_decoder.json	870	CF benchmark: all four models + paired bootstrap (Table 9)
gap_decode.json	1,479	gap-gene decoding (Table 10) + dosage-decodability CI
hb_boundaries.json	1,101	per-line Hb boundary positions
grn_real.json	2,769	exact attractor census, 8 conditions (Tables 3, 28)
grn_learning_curve.json	775	basin learning curve (Table 4)
grn_gnn.json	539	uniform-sampling degeneracy record
grn_gnn_balanced.json	250	balanced result + exact basin fraction 0.007725
grn_crosscheck.json	—	AEON/mpbn second-method attractor verification
grn_multicell.json	—	4-cell port: all-16 SLP prepatter scan (Table 7)
turing_sweep.json	3,011	14-point sweep (Table 1) + power-law fit
turing_cnn.json	592	surrogate result + retraction record
turing_sims_cache.json	255,116	171 cached dispersion curves and selections
turing_morphometrics.json	481	spot morphometrics + 96^2 cross-check
results.json	1,955	validation-suite numbers (analytic checks, Derrida, threshold readouts)

E Test-suite inventory

33 tests, all hermetic (no network, no optional data):

File	tests	Covers
test_bcd.py	5	gradient loading, fits, F -test, bootstrap
test_bnet.py	4	parser, synchronous semantics, attractors
test_gapgene.py	4	immunofluorescence handling, kNN decoder
test_grn.py	4	toggle switch, Derrida machinery
test_math_verify.py	4	symbolic verification of all derived formulas vs. committed numbers
test_morphogen.py	4	Turing solver, band, CFL, wavenumber
test_positional.py	4	threshold readout, error bounds
test_segment_polarity.py	5	reference model behaviour

F Claim-to-file traceability

Each headline number of this report, with the committed file and JSON key that produced it. “jq-path” means the value can be read with jq ‘<path>’ results/<file> on the committed repository state.

Table 30: Traceability of headline claims.

Claim	Value	jq path (file)
CNN CF-decoder RMSE	2.38% EL	<code>.models.cnn1d_profile_3seed.rmse_pctEL</code> (<code>cf_decoder.json</code>)
Threshold-model RMSE	4.94% EL	<code>.models.liu_threshold_trueD.rmse_pctEL</code> (<code>cf_decoder.json</code>)
Paired bootstrap gain CI	[1.99, 3.11]% EL	<code>.paired_bootstrap_cnn_minus_liu.ci95_pctEL</code> (<code>cf_decoder.json</code>)
DCLS Spearman ([0, 0.5])	$\rho = 0.356, p = 7.9 \times 10^{-16}$	<code>.windows."0.0-0.5".dcls_rho</code> / <code>.dcls_p</code> (<code>bcd_windows.json</code>)
DCLS OLS slope	0.0117, [0.0060, 0.0174]	CI <code>.slope_per_dose /</code> <code>.slope_ci95 (dcls_ols.json)</code>
λ on published window	0.189 EL	<code>.windows."0.2-0.8".lam_median</code> (<code>bcd_windows.json</code>)
Two-component rejection	0.74%	<code>.f_test_reject_fraction.p05</code> (<code>bcd_sdd.json</code>)
Dosage decodability	9.91 vs 6.21% EL	<code>.dosage_decodability.*</code> (<code>gap_decode.json</code>)
wg-ON basin fraction	0.7725%	<code>.wg_on_basin_fraction_uniform</code> (<code>grn_gnn_balanced.json</code>)
Learning-curve knee	$> 98\%$ at $n \approx 60$	<code>.curve</code> (<code>grn_learning_curve.json</code>)
Upshift range	14.7–58.9%	<code>.rows[] .dev</code> (tur- <code>ing_sweep.json</code>)
CNN surrogate (retracted)	0.292 vs 0.272	<code>.linear_plus_cnn_correction_relRMSE</code> vs <code>.linear_only_relRMSE</code> (<code>turing_cnn.json</code>)